




Introduction to Data Management

**** The “Flipped” Edition ****

Lecture #4

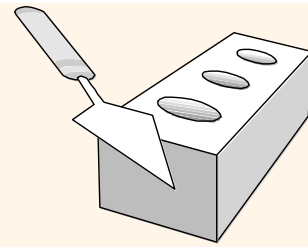
(ER Based DB Design)



Instructor: Mike Carey
mjcarey@ics.uci.edu




Today's Notices




- ❖ Continue watching the wiki page:
 - <http://www.ics.uci.edu/~cs122a/>
- ❖ And the Piazza page:
 - piazza.com/uci/fall2021/cs122aeecs116/home
- ❖ Partners: He/she is your “brainstorming buddy”!
 - Individual HW submissions (*not* team submissions)
 - See the partner-related part of the first HW assignment
- ❖ HW#1 is out! (Post questions on Piazza if needed)
 - *SWOOSH.com* (watch out, Zoom & Piazza!)
 - Please don't discuss *solutions* on Piazza!
- ❖ Quiz timing (for Quiz > 0)
 - Available from Wed 3PM – Fri **3PM** each week (**NEW PLAN!**)

Relational Database: Definitions



- ❖ Relational database: a set of *relations*
- ❖ Relation: consists of 2 parts:
 - *Instance*: a *table*, with rows and columns.
#rows = *cardinality*, #fields = *degree* or *arity*.
 - *Schema*: specifies name of relation, plus name and type of each column.
 - E.g, Students (*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).
- ❖ Can think of a relation as a *set* of rows or *tuples* (i.e., all rows are distinct) in the pure relational model (*vs.* reality of SQL 😊)

Example Instance of Students Relation



sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

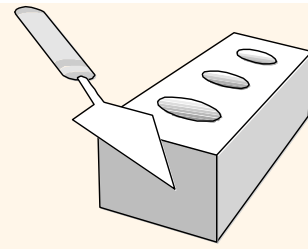
❖ Cardinality = 3, degree = 5, all rows distinct



Relational Query Languages

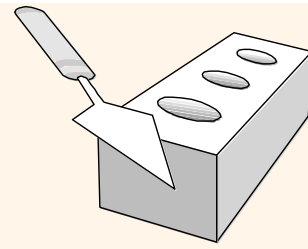
- ❖ A major strength of the relational model: supports simple, powerful *querying* of data.
- ❖ Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
 - The key: precise (and set-based) semantics for relational queries.
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

SQL Query Language (Preview)



- ❖ Developed by IBM (System R) in the 1970s
- ❖ Need for a standard, since it is used by many vendors (Oracle, IBM, Microsoft, ...)
- ❖ ANSI/ISO Standards:
 - SQL-86
 - SQL-89 (minor revision)
 - SQL-92 (major revision, very widely supported)
 - SQL-99 (major extensions, current standard)

SQL Query Language (Preview)



❖ To find all 18-year-old students, we can write:


```
SELECT *  
FROM Students S  
WHERE S.age=18
```

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

Querying Multiple Relations



- ❖ What does the following query compute?

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade="A"
```

Given the following instances of Students and Enrolled:

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

We will get:

S.name	E.cid
Smith	Topology112



Integrity Constraints (ICs)

- ❖ **IC:** Condition that must be true for *any* instance of the database; e.g., *domain constraints*.

ICs are specified when schema is defined.

ICs are checked when relations are modified.

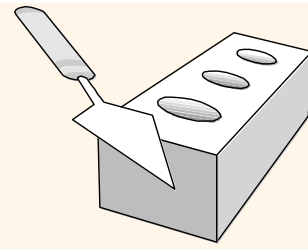
- ❖ A *legal* instance of a relation is one that satisfies all specified ICs.

DBMS should not allow illegal instances.

- ❖ If the DBMS checks ICs, stored data is more faithful to real-world meaning.

Avoids data entry errors (centrally), too!

Primary Key Constraints



❖ A set of fields is a key for a relation if :



1. No two distinct tuples can have the same values in all key fields, and

2. This is not true for any subset of the key.

Part 2 false? In that case, this is a “*superkey*”.


If there's > 1 key for a relation, one of the keys is chosen (by DBA) to be the *primary key*.

The others are referred to as *candidate keys*.



❖ E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.

Primary and Candidate Keys in SQL



❖ Possibly several candidate keys (specified using **UNIQUE**), but one is chosen as the *primary key*.


❖ Used carelessly, an IC can prevent the storage of database instances that arise in practice!

❖ “For a given student + course, there is a single grade.” **vs.**
“Students can take only one course and receive a single grade for that course; further, no two students in a course may ever receive the same grade.”

```
CREATE TABLE Enrolled  
(sid VARCHAR(20)  
  cid VARCHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid,cid) )
```

```
CREATE TABLE Enrolled  
(sid VARCHAR(20)  
  cid VARCHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid),  
  UNIQUE (cid, grade) ),11
```

Foreign Keys, Referential Integrity



- ❖ Foreign key: Set of fields in one relation used to “refer” to a tuple in another relation. (Must refer to the primary key of the other relation.) Like a “logical pointer”.
- ❖ E.g., *sid* is a foreign key referring to **Students**:
 - Enrolled(*sid*: string, *cid*: string, *grade*: string)
 - If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.



Foreign Keys in SQL

- ❖ Ex: Only students listed in the Students relation should be allowed to enroll for courses.

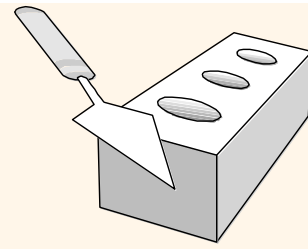
```
CREATE TABLE Enrolled  
  (sid VARCHAR(20), cid VARCHAR(20), grade CHAR(2),  
   PRIMARY KEY (sid, cid),  
   FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

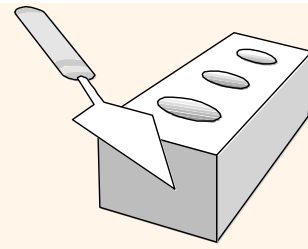
Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8



Enforcing Referential Integrity

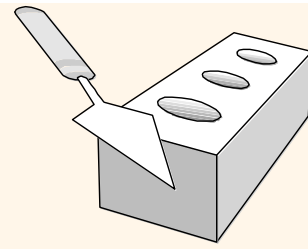
- ❖ Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- ❖ What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- ❖ What should be done if a Students tuple is deleted?
 - Also delete all Enrolled tuples that refer to it. Or...
 - Disallow deletion of a Students tuple if it is referred to.
 - Set *sid* in Enrolled tuples that refer to it to a *default sid*.
 - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting '*unknown*' or '*inapplicable*'.)
- ❖ Similar if primary key of Students tuple is updated.



Referential Integrity in SQL

- ❖ SQL/92 and SQL:1999 support all 4 options on deletes and updates.
 - Default is **NO ACTION** (*delete/update is rejected*)
 - CASCADE** (also delete all tuples that refer to the being-deleted tuple)
 - SET NULL / SET DEFAULT** (sets foreign key value of the referring tuples)


```
CREATE TABLE Enrolled
(sid VARCHAR(20),
cid VARCHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```




Where Do ICs Come From?

- ❖ ICs are based upon the semantics of the real-world enterprise that is being described in the database relations (perhaps via an E-R schema)
- ❖ We can check a database instance to see if an IC is violated, but we can **NEVER** infer that an IC is true by looking at an instance.
 - An IC is a statement about *all possible* instances!
 - For example, we know *name* is not a key, but the assertion that *sid* is a key is given to us.
- ❖ Key and foreign key ICs are the most common; more general ICs supported too.

Logical DB Design: ER to Relational



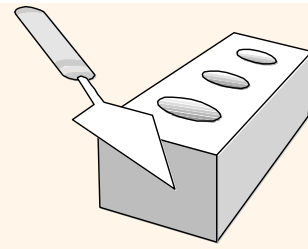
❖ Entity sets to tables:



```
CREATE TABLE Employees  
(ssn CHAR(11),  
name VARCHAR(20),  
lot INTEGER,  
PRIMARY KEY (ssn))
```

❖ Notes:

- PRIMARY KEYs are NOT NULL by default
- All other fields are NULL(-able) by default
- One can say NOT NULL

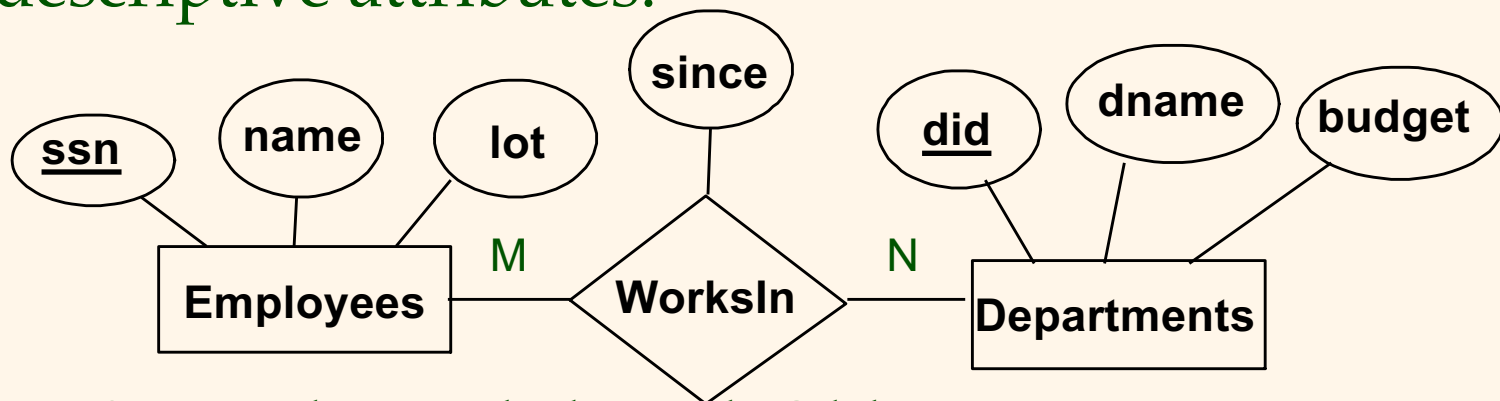


Relationship Sets to Tables


- ❖ In translating a relationship set to a relation, attributes of the relation should include:
 - Keys for each participating entity set (as foreign keys).
 - This set of attributes forms a *superkey* for the relation.

All descriptive attributes.



```
CREATE TABLE Works_In(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (ssn, did),  
  FOREIGN KEY (ssn)  
    REFERENCES Employees,  
  FOREIGN KEY (did)  
    REFERENCES Departments)
```




Key Constraints (Review)




- ❖ Each dept has at most one manager, according to the key constraint on Manages.




1-to-1



1-to Many




Many-to-1



Many-to-Many

Translation to relational model?



Translating ER Diagrams with Key Constraints

- ❖ Map the relationship to a table (Manages):
 - Note that **did** (alone) is the key!
 - Still separate tables for Employees and Departments.
- ❖ But, since *each* department has a *unique* manager, we could choose to fold Manages right into Departments.
(Q: Why do this...?)

```
CREATE TABLE Manages (  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  FOREIGN KEY (did) REFERENCES Departments)
```



vs.

```
CREATE TABLE Departments2 (  
  did INTEGER,  
  dname VARCHAR(20),  
  budget REAL,  
  mgr_ssn CHAR(11),  
  mgr_since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (mgr_ssn) REFERENCES Employees)
```



Note: The relationship info has been pushed to the N-side's entity table!

Properly Reflecting Key Constraints



❖ So what are the translated relationship table's keys (*etc.*) when...

FooBar is M:N? → FooBar(fooId, barId, baz)

FooBar is N:1? → FooBar(fooId, barId, baz)

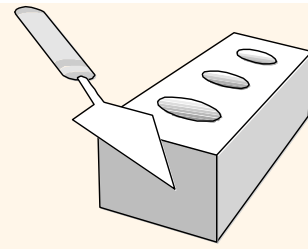
FooBar is 1:N? → FooBar(fooId, barId, baz)

FooBar is 1:1? → FooBar(fooId, barId, baz) (Note: unique)

Ensures unique
Foo/Bar pairs

Ensures one Bar
per Foo entity

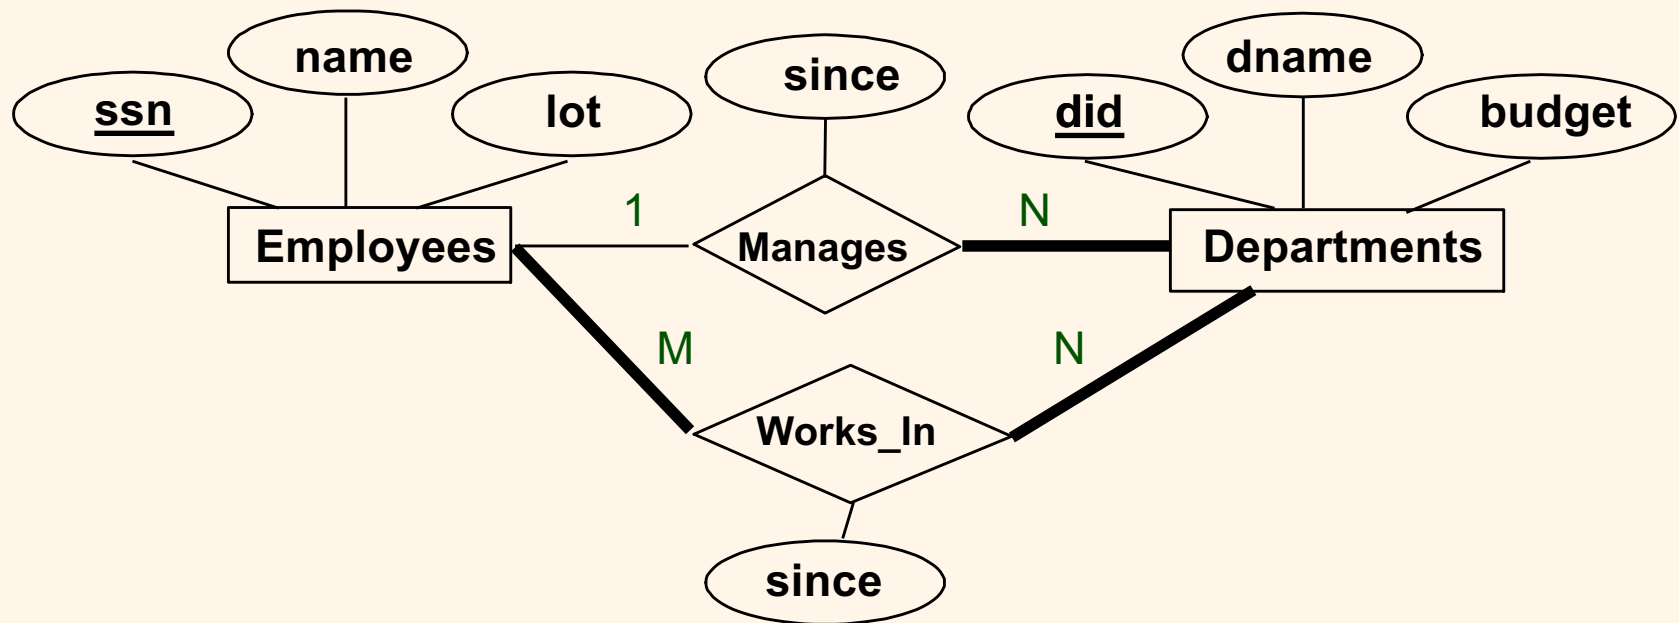
Review: Participation Constraints




❖ Does every department have a manager?

If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).

- Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!!)



Participation Constraints in SQL



- ❖ We can capture participation constraints involving the *N-side* entity set in a binary relationship, but little else (without resorting to the use of *triggers*).

```
CREATE TABLE Department2 (  
    did INTEGER,  
    dname VARCHAR(20),  
    budget REAL,  
    mgr_ssn CHAR(11) NOT NULL,  
    mgr_since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (mgr_ssn) REFERENCES Employees,  
    ON DELETE NO ACTION*) (*or: RESTRICT)
```

To Be Continued...

