

Introduction to Data Management

**** The “Flipped” Edition ****


Lecture #5

(ER \rightarrow Relational Mapping)


Instructor: Mike Carey
mjcarey@ics.uci.edu



Today's Notices





- ❖ Keep watching the wiki page:
 - <http://www.ics.uci.edu/~cs122a/>
- ❖ And, of course, follow the Piazza page:
 - piazza.com/uci/fall2021/cs122aeecs116/home
- ❖ HW#1 is in flight! (Keep Q's coming on Piazza...)
 - *SWOOSH.com*
 - And, of course, avoid discussing *solutions* publicly here!
- ❖ Remember that quizzes run from **Wed 3PM – Fri 3PM**
 - Don't miss the free points! (You will if you fall behind!)



Review: Weak Entities

- ❖ A *weak entity* can be identified (uniquely) only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.





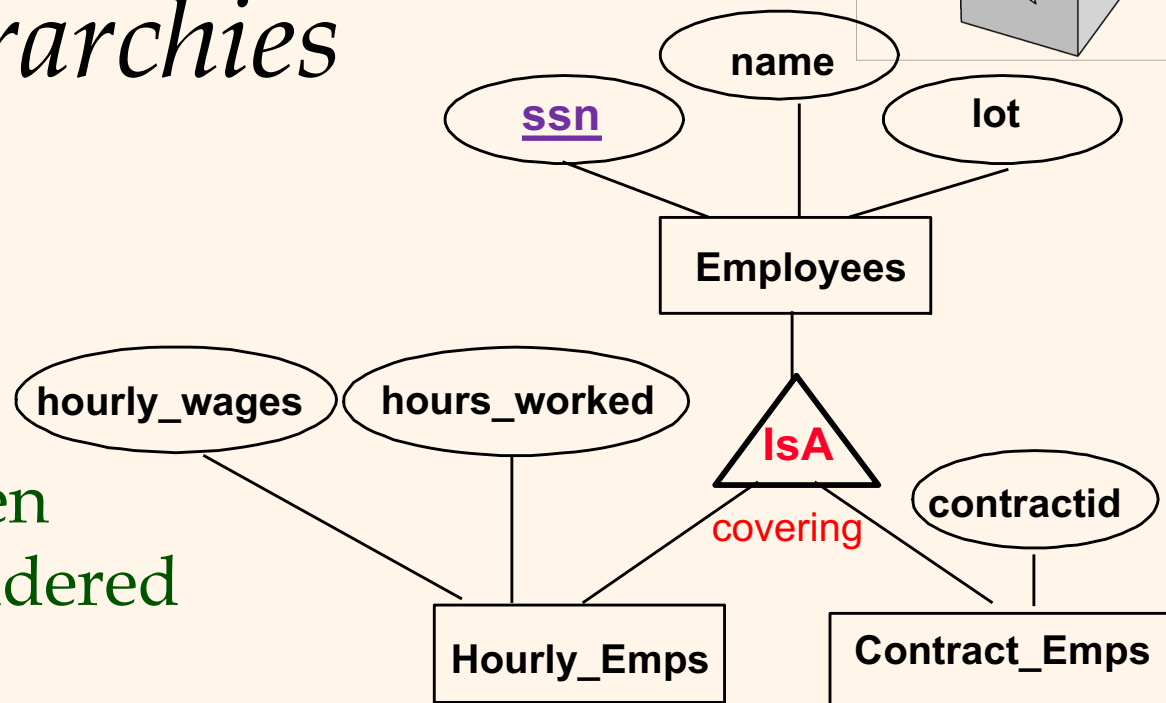
Translating Weak Entity Sets

- ❖ Weak entity set and identifying relationship set are translated into a *single table*.
 - When the owner entity is deleted, all of its owned weak entities must also be deleted.


```
CREATE TABLE Dependents2 (  
    dname VARCHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (dname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees  
    ON DELETE CASCADE)
```

Review: IsA Hierarchies

- ❖ As in C++, or other PLs, attributes are inherited.
- ❖ If we declare A **IsA** B, then every A entity is also considered to be a B entity.




- ❖ *Overlap constraints*: Can employee Joe be an Hourly_Emps as well as a Contract_Emps entity? (*disjoint* if not)
- ❖ *Covering constraints*: Must every Employees entity be either an Hourly_Emps or a Contract_Emps entity? (*covering* if so)



From IsA Hierarchies to Relations

- ❖ **Most general and “clean” approach** (recommended):
 - 3 relations: Employees, Hourly_Emps, & Contract_Emps.
 - *Hourly_Emps*: Every employee recorded in Employees. For hourly emps, *extra* info recorded in Hourly_Emps (*hourly_wages*, *hours_worked*, *ssn*); delete Hourly_Emps tuple if its referenced Employees tuple is deleted.
 - Queries about *all* employees easy; those involving *just* Hourly_Emps require a join to access the extra attributes.
- ❖ **Another alternative: Hourly_Emps & Contract_Emps.**
 - *Ex: Hourly_Emps*(*ssn*, *name*, *lot*, *hourly_wages*, *hours_worked*)
 - If each employee must be in one of the two subclasses...
(Q: Can we *always* do this, then? A: Not w/o redundancy!)



IsA Hierarchy Translation Options

❖ I. “Delta table” approach (recommended):

- *Emps*(*ssn*, *name*, *lot*) (All Emps partly reside here)
- *Hourly_Emps*(*ssn**, *wages*, *hrs_worked*)
- *Contract_Emps*(*ssn**, *contractid*)

Things to consider:

- Expected queries?
- PK/unique constraints?
- Relationships/FKs?
- Overlap constraints?
- Space/time tradeoffs?


❖ II. “Union of tables” approach:

- *Emps*(*ssn*, *name*, *lot*)
- *Hourly_Emps*(*ssn*, *name*, *lot*, *wages*, *hrs_worked*)
- *Contract_Emps*(*ssn*, *name*, *lot*, *contractid*)

❖ III. “Mashup table” approach:

- *Emps*(*kind*, *ssn*, *name*, *lot*, *wages*, *hrs_worked*, *contractid*)

(**ssn* here is both a local PK as well as an FK referencing Emps)



IsA Considerations (cont'd.)

❖ Query convenience

- *Ex:* List the names of all Emps in lot 12A

❖ PK enforcement

- *Ex:* Make sure that ssn is unique for all Emps

❖ Relationship (FK) targets


- *Ex:* Lawyers table REFERENCES Contract_Emps

❖ Handling of overlap constraints

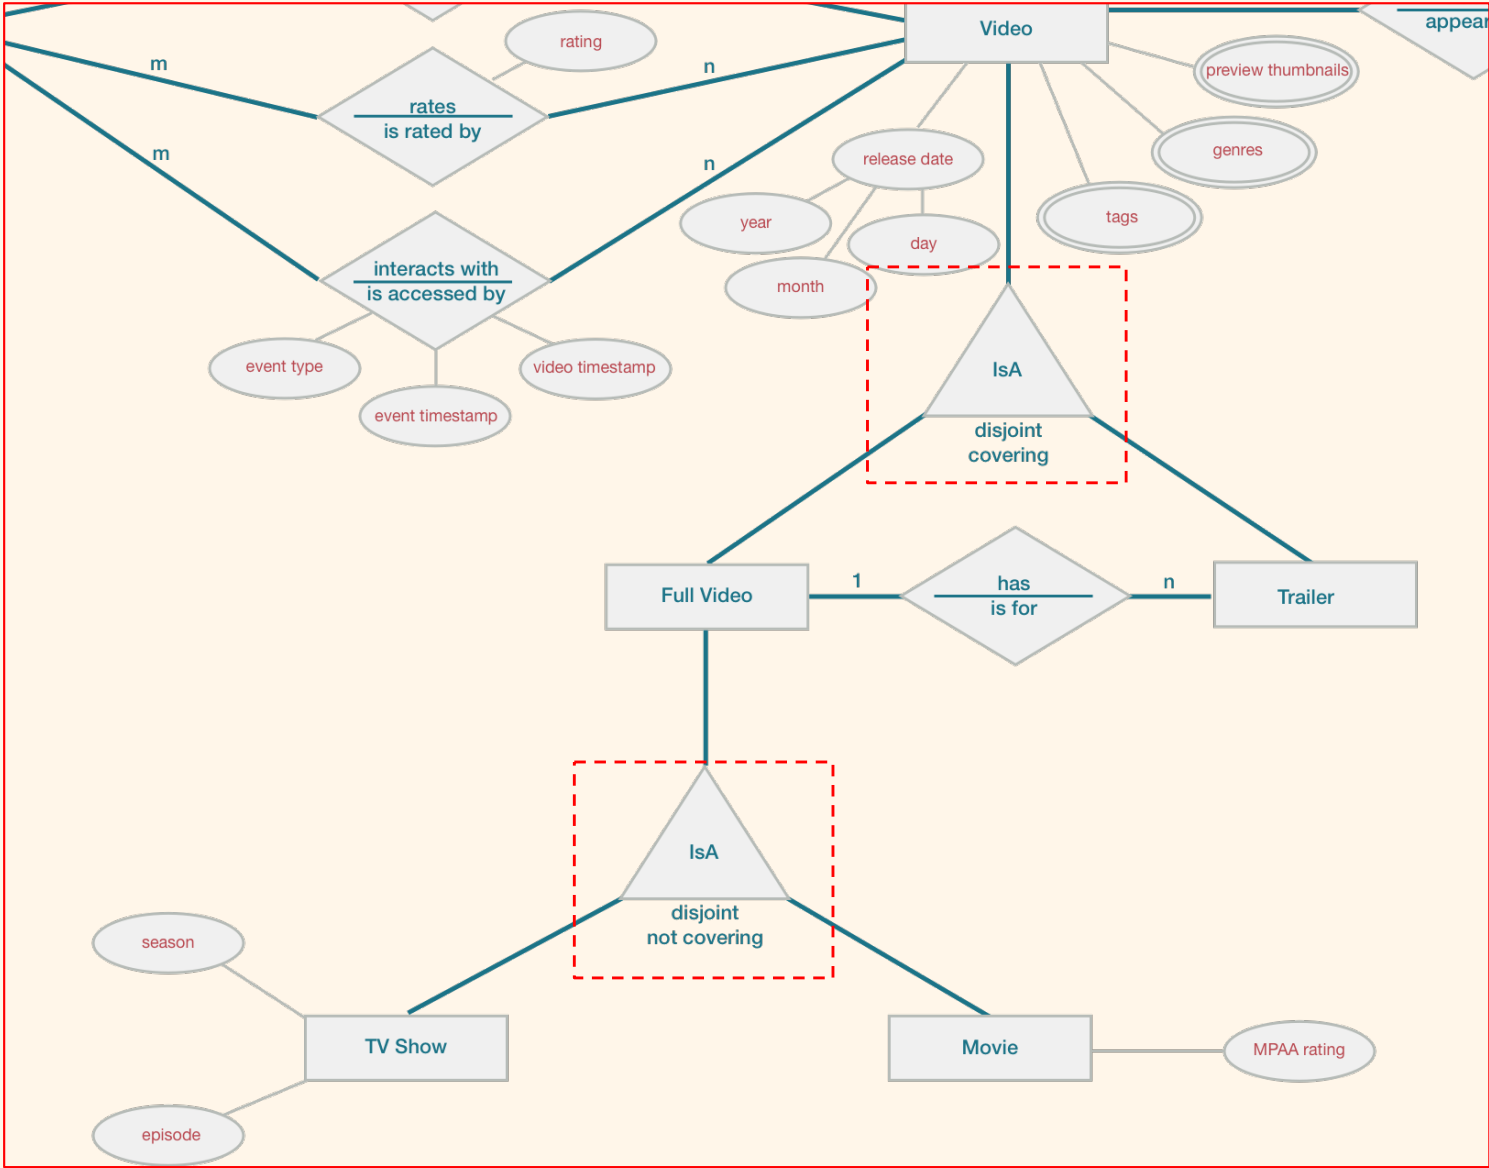
- *Ex:* Sally is under a contract for her hourly work


❖ Space and query performance tradeoffs

- *Ex:* List all the info about hourly employee 123
- *Ex:* What if most employees are “just plain employees”?



Logical Design for SQL (Ex. 2)

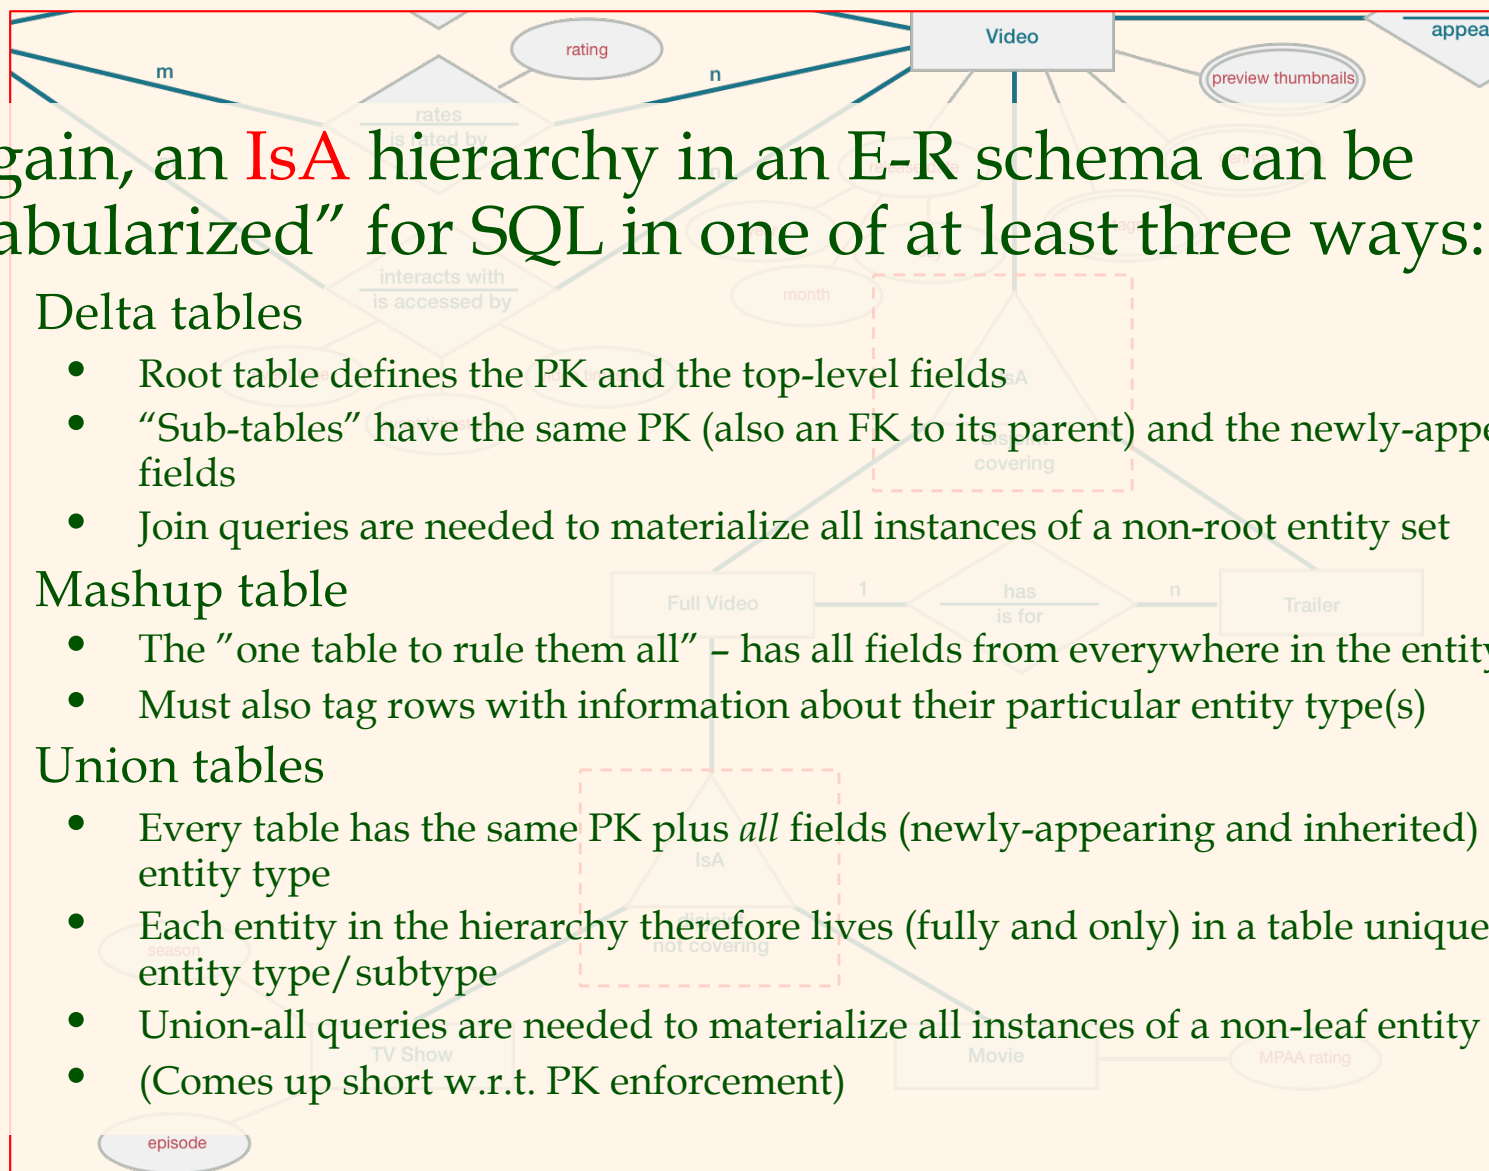





Logical Design for SQL (Ex. 2)

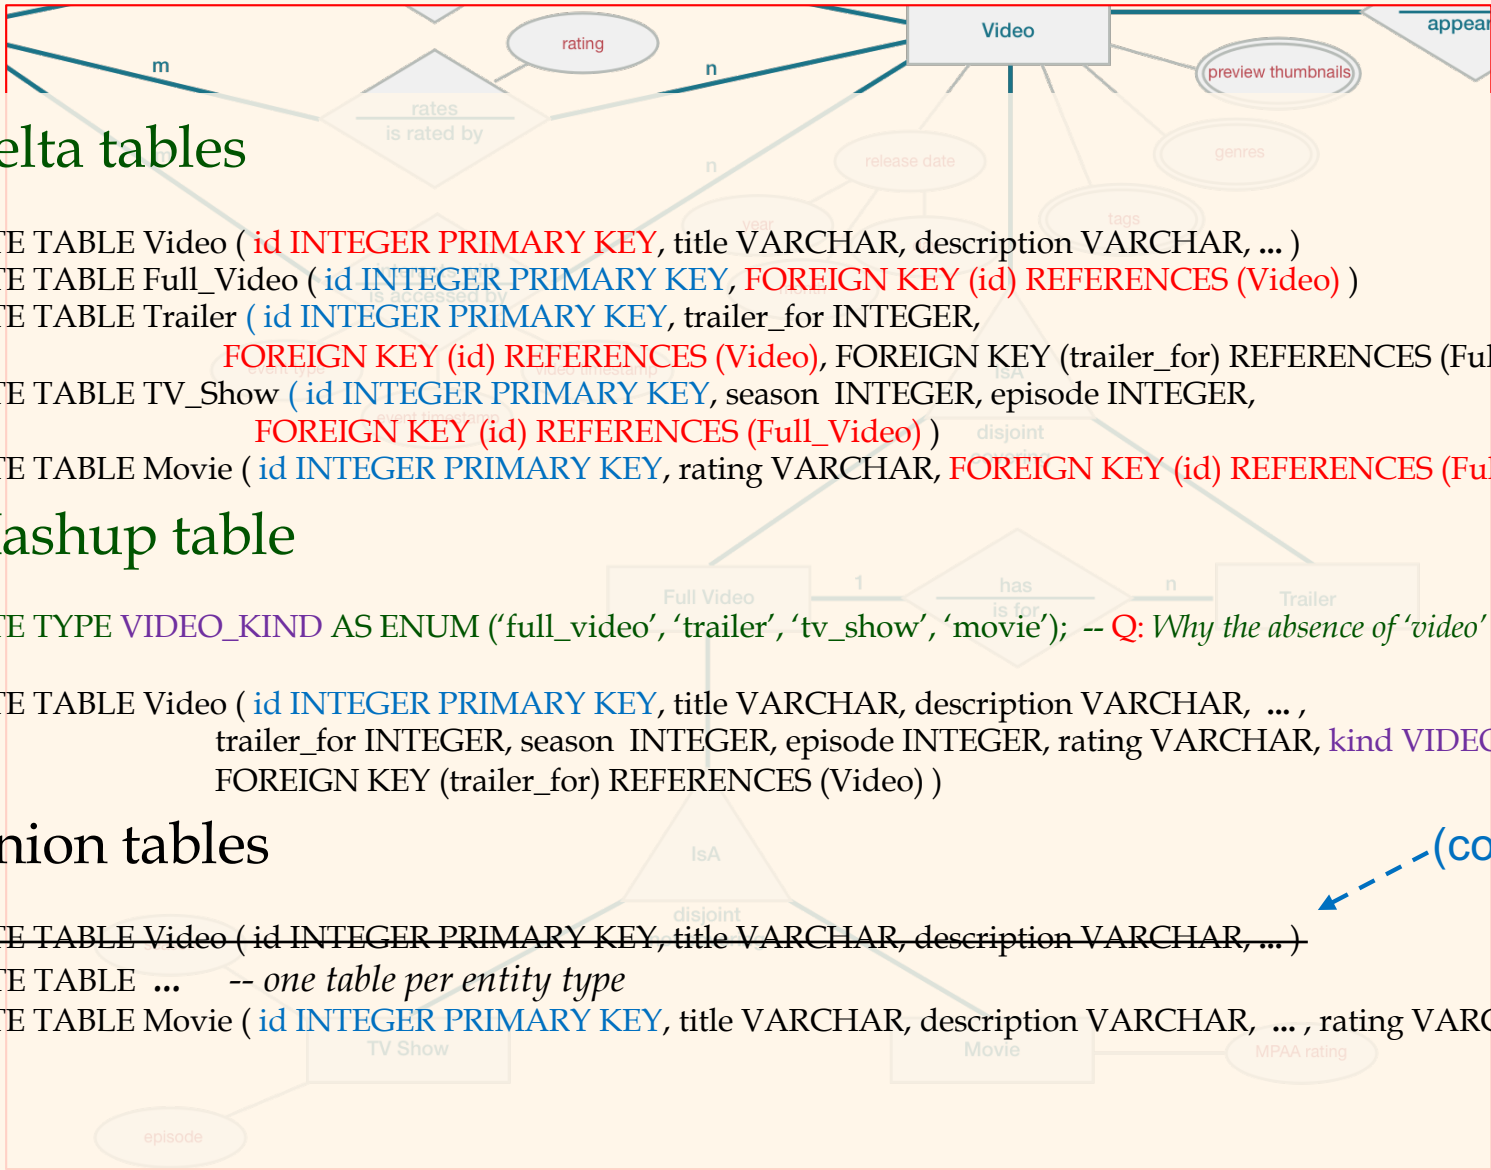
❖ Again, an **IsA** hierarchy in an E-R schema can be “tabularized” for SQL in one of at least three ways:

- Delta tables
 - Root table defines the PK and the top-level fields
 - “Sub-tables” have the same PK (also an FK to its parent) and the newly-appearing fields
 - Join queries are needed to materialize all instances of a non-root entity set
- Mashup table
 - The “one table to rule them all” – has all fields from everywhere in the entity hierarchy
 - Must also tag rows with information about their particular entity type(s)
- Union tables
 - Every table has the same PK plus *all* fields (newly-appearing and inherited) for its entity type
 - Each entity in the hierarchy therefore lives (fully and only) in a table unique to its entity type/subtype
 - Union-all queries are needed to materialize all instances of a non-leaf entity type
 - (Comes up short w.r.t. PK enforcement)





Logical Design for SQL (Ex. 2)



❖ Delta tables

```
CREATE TABLE Video ( id INTEGER PRIMARY KEY, title VARCHAR, description VARCHAR, ... )
CREATE TABLE Full_Video ( id INTEGER PRIMARY KEY, FOREIGN KEY (id) REFERENCES (Video) )
CREATE TABLE Trailer ( id INTEGER PRIMARY KEY, trailer_for INTEGER,
FOREIGN KEY (id) REFERENCES (Video), FOREIGN KEY (trailer_for) REFERENCES (Full_Video) )
CREATE TABLE TV_Show ( id INTEGER PRIMARY KEY, season INTEGER, episode INTEGER,
FOREIGN KEY (id) REFERENCES (Full_Video) )
CREATE TABLE Movie ( id INTEGER PRIMARY KEY, rating VARCHAR, FOREIGN KEY (id) REFERENCES (Full_Video) )
```

❖ Mashup table


```
CREATE TYPE VIDEO_KIND AS ENUM ('full_video', 'trailer', 'tv_show', 'movie'); -- Q: Why the absence of 'video' here?

CREATE TABLE Video ( id INTEGER PRIMARY KEY, title VARCHAR, description VARCHAR, ... ,
trailer_for INTEGER, season INTEGER, episode INTEGER, rating VARCHAR, kind VIDEO_KIND,
FOREIGN KEY (trailer_for) REFERENCES (Video) )
```

❖ Union tables


```
CREATE TABLE Video ( id INTEGER PRIMARY KEY, title VARCHAR, description VARCHAR, ... )
CREATE TABLE ... -- one table per entity type
CREATE TABLE Movie ( id INTEGER PRIMARY KEY, title VARCHAR, description VARCHAR, ... , rating VARCHAR)
```

(covering)



Mapping Advanced ER Features

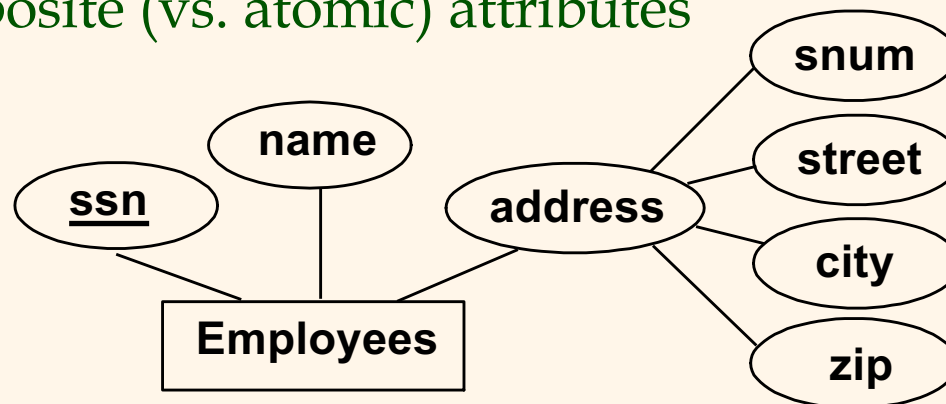
❖ Multi-valued (vs. single-valued) attributes




Employees_phones(ssn, phone)

- ssn is an FK in this table
- (ssn, phone) is its PK

❖ Composite (vs. atomic) attributes




Employees(ssn, name, address_snum, address_street, address_city, address_zip)



Mapping Advanced ER Features (cont.)


❖ Mandatory (vs. optional) attributes




Employees(ssn, name, age, gender)

- ssn is the PK
- name VARCHAR(20) **NOT NULL**
- Note: **PRIMARY KEY** → **NOT NULL**

❖ Derived (vs. stored) attributes






SQL Views (and Security)

- ❖ A view is just a relation, but we store its *definition* rather than storing the (materialized) set of tuples.

```
CREATE VIEW YoungStudents (name, login)
AS SELECT S.name, S.login
FROM Students S
WHERE S.age < 19;
```

- ❖ Views can be used to present needed information while hiding details of underlying table(s).
 - Given YoungStudents (but not Students), we can see (young) students *S* with only their names and logins.



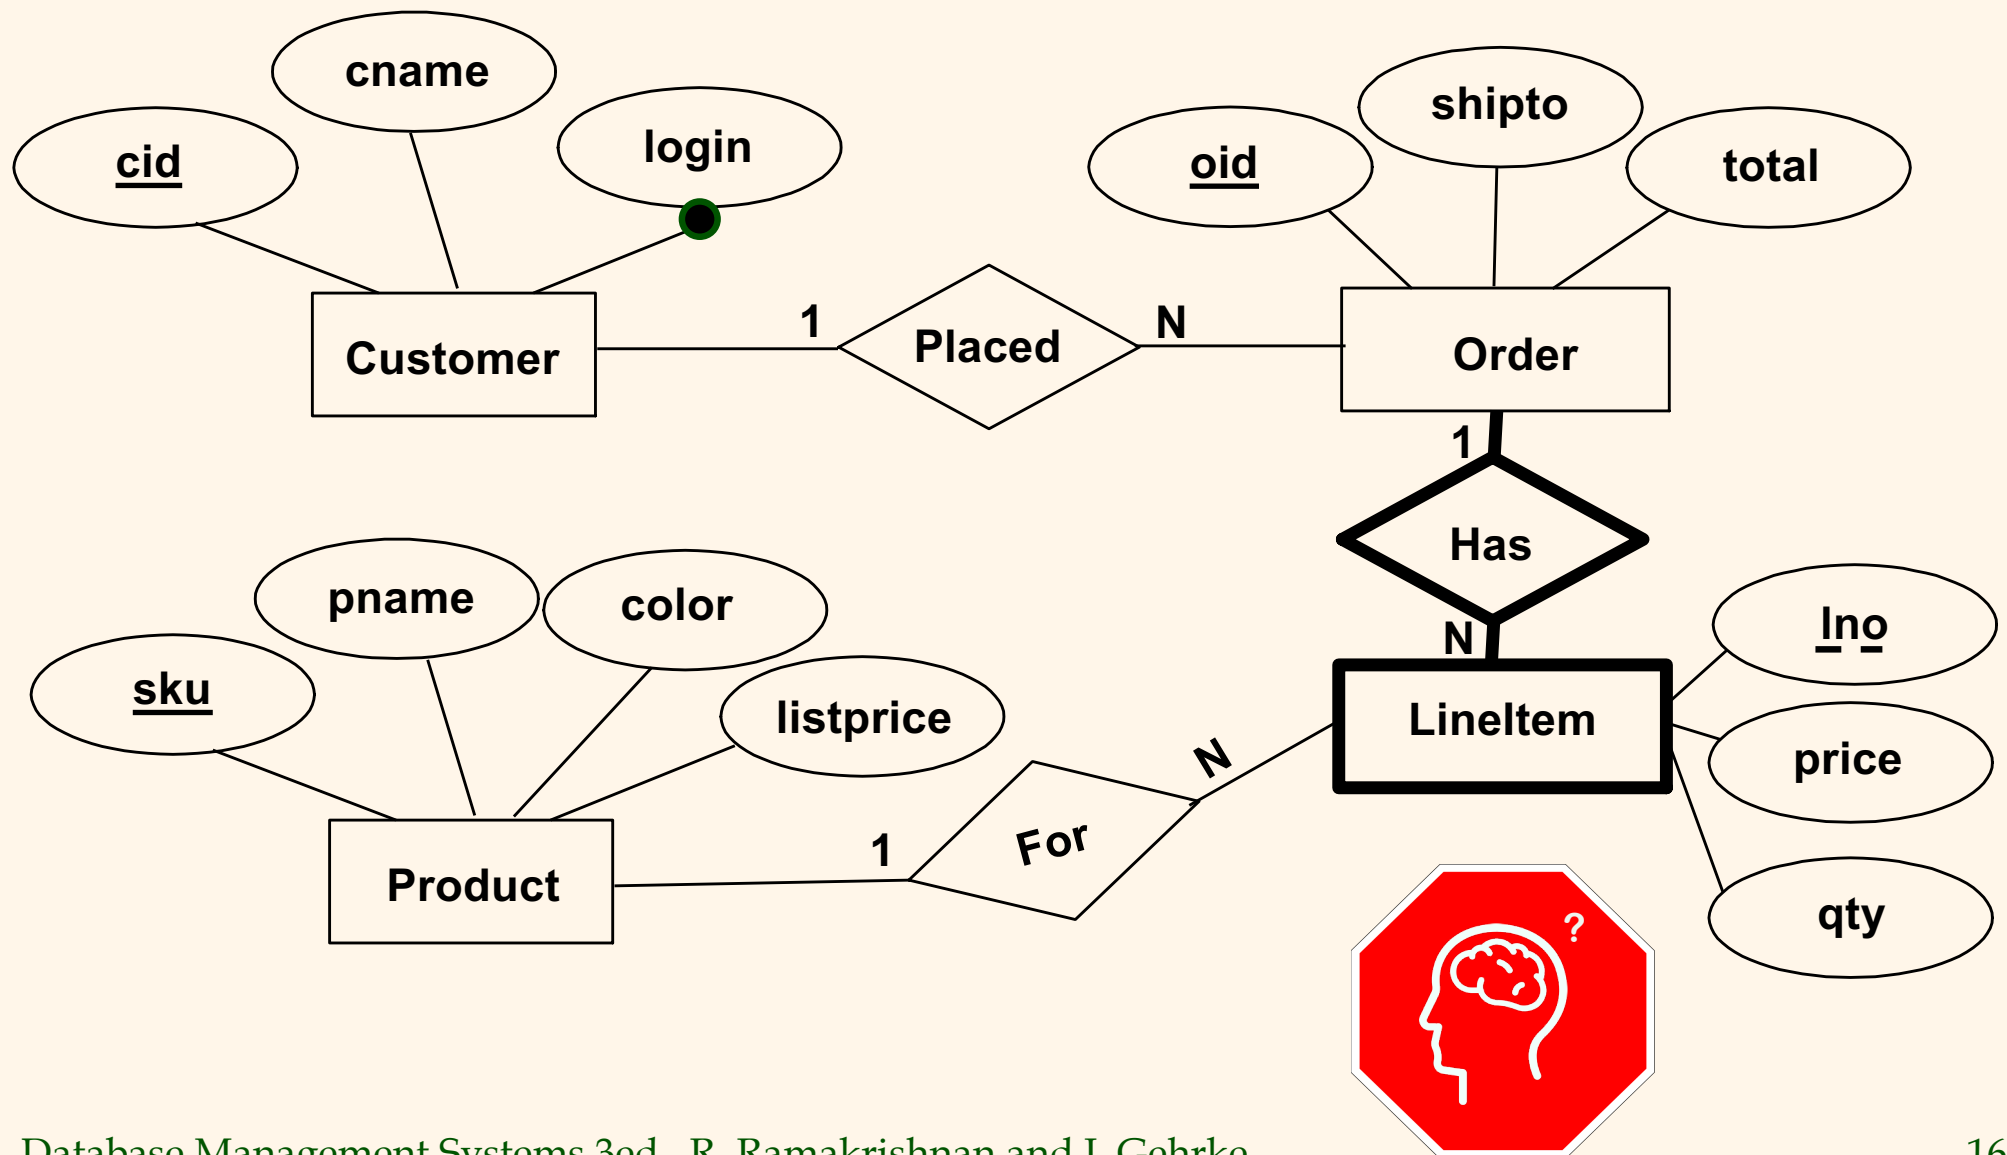

SQL Views (Cont'd.)

❖ Other view uses in our ER translation context might include:


- *Derived attributes, e.g., age (vs. birthdate)*
- Simplifying/eliminating join paths (for SQL)
- Beautifying the “Mashup table” approach (to IsA)

```
CREATE VIEW EmployeeView (ssn, name, bdate, age)
AS SELECT E.ssn, E.name, E.bdate,
        date_part('year', age(E.bdate))::int
FROM Employees E;
```

Review: Putting the Basics Together




Review: Putting It Together (Cont'd.)



```
CREATE TABLE Customer (  
  cid INTEGER,  
  cname VARCHAR(50),  
  login VARCHAR(20)  
    NOT NULL,  
  PRIMARY KEY (cid),  
  UNIQUE (login))  
  
CREATE TABLE Product (  
  sku INTEGER,  
  pname VARCHAR(100),  
  color VARCHAR(20),  
  listprice DECIMAL(8,2),  
  PRIMARY KEY (sku))  
  
CREATE TABLE Order (  
  oid INTEGER,  
  custid INTEGER,  
  shipto VARCHAR(200),  
  total DECIMAL(8,2),  
  PRIMARY KEY (oid),  
  FOREIGN KEY (custid) REFERENCES Customer))  
  
CREATE TABLE LineItem (  
  oid INTEGER,  
  lno INTEGER,  
  price DECIMAL(8,2),  
  qty INTEGER,  
  sku INTEGER,  
  PRIMARY KEY (oid, lno),  
  FOREIGN KEY (oid) REFERENCES Order  
    ON DELETE CASCADE),  
  FOREIGN KEY (sku) REFERENCES Product))
```

Review: Putting It Together (Cont'd.)



Customer

cid	cname	login
1	Smith, James	jsmith@aol.com
2	White, Susan	<u>suzie@gmail.com</u>
3	Smith, James	js@hotmail.com

Product

sku	pname	color	listprice
123	Frozen DVD	null	24.95
456	Graco Twin Stroller	green	199.99
789	Moen Kitchen Sink	black	350.00


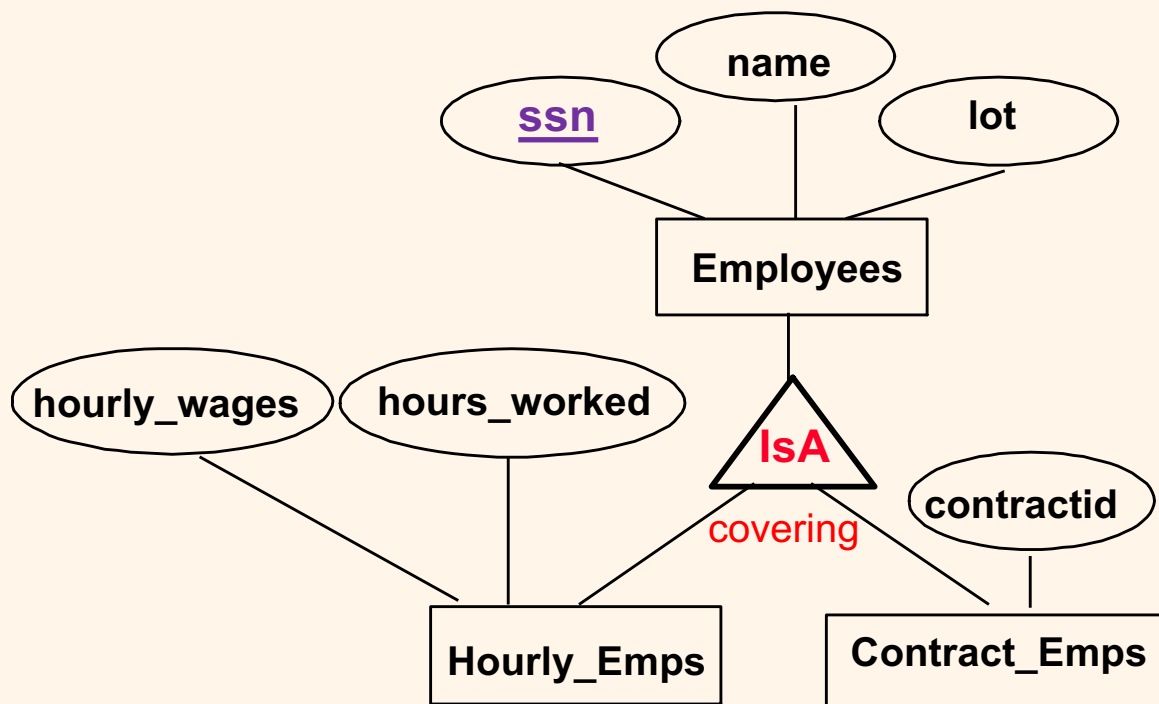

Order


oid	custid	shipto	total
1	3	J. Smith, 1 Main St., USA	199.95
2	1	Mrs. Smith, 3 State St., USA	300.00

LineItem

oid	lno	price	qty	item
1	1	169.95	1	456
1	2	15.00	2	123
2	1	300.00	1	789

Wait! Clarifying IsA Mappings...

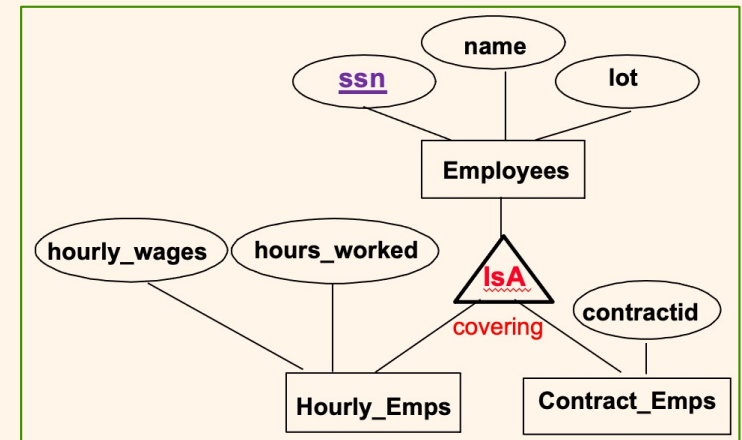




Delta Table Approach

Employees

ssn	name	lot
1	Joe	12
2	Shelly	16
3	Arvind	null
4	Chen	11



Note 1: Joe can't really exist here (*covering*)

Note 2: Chen couldn't exist if we'd said *disjoint*


Hourly_Emps

ssn	wages	hrs_worked
2	15	40
4	250	10

Contract_Emps

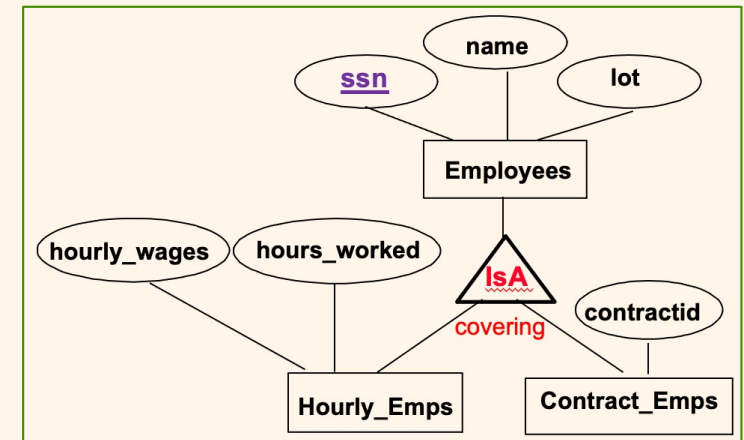
ssn	contract_id
3	101
4	102

Union Table Approach



Employees

ssn	name	lot
4	Joe	12



Note 1: Joe can't really exist here (*covering*)

Hourly_Emps



ssn	name	lot	wages	hrs_worked
2	Shelly	16	15	40
4	Chen	11	250	10

Contract_Emps

ssn	name	lot	contract_id
3	Arvind	<i>null</i>	101
4	Chen	11	102

Note 3: No CREATE TABLE way to prevent *ssn* duplication

Mashup Table Approach



Employees


ssn	is_hourly	is_contract	name	lot	wages	hrs_worked	contract_id
1	false	false	Joe	12	null	null	null
2	true	false	Shelly	16	15	40	null
3	false	true	Arvind	null	null	null	101
4	true	true	Chen	11	250	10	102

Note 1: Joe can't really exist here (*covering*)

Note 4: Might want to create VIEWS on top...

Relational Model and E-R Schema

Translation: Summary



- ❖ Relational model: Tabular data representation.
- ❖ Simple and intuitive, *very* widely used (RDBMSs)
- ❖ Integrity constraints – specified by DBA *based on application semantics*. (DBMS prevents violations)
 - Most important ICs: Primary and foreign keys (PKs and FKs)
 - In addition, we also have domain (column type) constraints
- ❖ High-level query languages (including SQL!)
- ❖ Rules to translate E-R to relational model
 - Can be done by a human or automatically (via a tool)
 - Entities, relationships, attributes; cardinality, participation,...
 - IsA handling; composite, multi-valued, and derived attributes