




Introduction to Data Management

**** The “Flipped” Edition ****

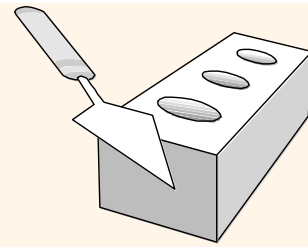
Lecture #9

(Relational Languages I)

Instructor: Mike Carey
mjcarey@ics.uci.edu




Today's Notices



- ❖ Keep one eye on the wiki page...
 - <http://www.ics.uci.edu/~cs122a/>
- ❖ ... and the other eye on Piazza Q&A!
 - piazza.com/uci/fall2021/cs122aeecs116
- ❖ HW #2 should be done!
 - Due “today” at 6PM to be on time
 - Then comes the 24-hour grace period (-10%)
- ❖ Today's lecture plan:
 - Relational languages – *the next frontier...*
 - **Note:** Today's material won't be on Midterm 1! (☺)



Quick Time Check...




Topic Coverage and Exam Schedule

Syllabus

Topic	Reading (Required!)
Databases and DB Systems	Ch. 1
Entity-Relationship (E-R) Data Model	Ch. 6.1-6.5, 6.8-6.9
Relational Data Model	Ch. 2.1-2.4, 3.1-3.2
E-R to Relational Translation	Ch. 6.6-6.7
Relational Design Theory	Ch. 7.1-7.4.2
Midterm Exam 1	Fri, Oct 22 (during lecture time)
Relational Algebra	Ch. 2.5-2.7
Relational Calculus	⇒ Wikipedia: Tuple relational calculus
SQL Basics (SPJ and Nested Queries)	Ch. 3.3-3.5
SQL Analytics: Aggregation, Nulls, and Outer Joins	Ch. 3.6-3.9, 4.1
Advanced SQL: Constraints, Triggers, Views, and Security	Ch. 4.2, 4.4-4.5, 4.7
Midterm Exam 2	Mon, Nov 15 (during lecture time)
Storage	Ch. 12.1-12.4, 12.6-12.7
Indexing	Ch. 14.1-14.4, 14.5
Physical DB Design	Ch. 14.6-14.7, 15.1-15.3, 15.5.3
Semistructured Data Management (<i>a.k.a.</i> NoSQL)	Ch. 8.1, ⇒ AsterixDB SQL++ Primer , ⇒ Couchbase SQL++ Book
Data Science 1: Advanced SQL Analytics	Ch. 5.5, 11.3
Data Science 2: Notebooks, Dataframes, and Python/Pandas	Lecture notes and Jupyter notebook
Basics of Transactions	Ch. 4.3, Ch. 17
Endterm Exam	Fri, Dec 3 (during lecture time)


Midterm Exam 1

Time: Fri, Oct 22, Lecture Time
Place: SSLH 100





First: A Word on Learning...

❖ "When **I** was your age..." (☹)



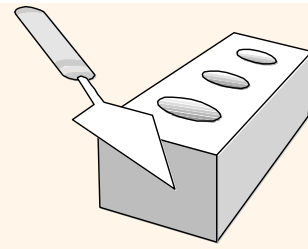
■ No web! (Just the





First: A Word on Learning...


- ❖ "When **I** was your age..." (☺)
 - 8-bit μ processors, LEDs, hex keypad, 16-bit for its address space, ...
 - FORTRAN, Pascal, C, Lisp, Snobol, APL, Quel ...
 - Unix first appeared, Unix/INGRES DB 64K process design point, and 1MB of memory was amazing!
 - No web! (Just the early ArpaNet)
- ❖ **Fast forward 35 years...**
 - Your cell phones dwarf our ~1980 computing platforms
 - Python, Java, Go, C++, Ruby, SQL, SQL++, ...
 - Twitter... (☹)
 - **WHAT THIS MEANS:** It's critical to *learn* how to read and how to *learn*, how to *search* for info/resources, etc...!



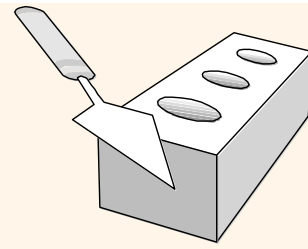
On to Relational Query Languages!

- ❖ Query languages: Allow manipulation and *retrieval* of data from a database.
- ❖ Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- ❖ Query Languages *≠* programming languages!
 - QLs not expected to be “Turing complete.”
 - QLs not intended to be used for complex calculations.
 - *QLs support easy, efficient access to large data sets.*

Formal Relational Query Languages



- ❖ Two mathematical Query Languages form the basis for “real” languages (e.g., SQL), and for their implementation:
 - Relational Algebra: More **operational**, very useful for representing execution plans.
 - Relational Calculus: Lets users describe **what** they want, rather than how to compute it.
(**Non-operational**, or declarative.)




Preliminaries

- ❖ A query is applied to *relation instances*, and the result of a query is *also* a relation instance.
 - *Schemas of input* relations for a query are *fixed* (but query will run regardless of instance!)
 - The *schema for the result* of a given query is also *fixed!* Determined by applying the definitions of the query language's constructs.
- ❖ Positional vs. named-field notation:
 - Positional notation easier for formal definitions, but named-field notation far more readable.
 - Both used in SQL (but try to avoid positional stuff!)

Example Instances

- ❖ “Sailors” and “Reserves” relations for our examples.
- ❖ We’ll use positional or named field notation and assume that names of fields in query results are “inherited” from names of fields in query input relations (when possible).

R1



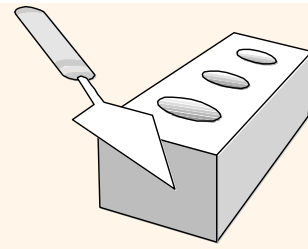
<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



Relational Algebra

❖ Basic operations:

- Selection (σ) Selects a subset of rows from relation.
- Projection (π) Omits unwanted columns from relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.

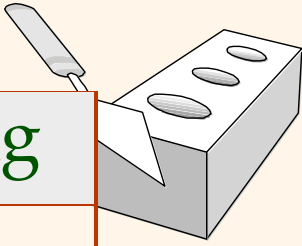
❖ Additional operations:

- Intersection, join, division, renaming: Not essential, but (very!) useful. (I.e., don't add expressive power, but...)

❖ Since each operation returns a relation, **operations can be composed!** (Algebra is “closed”.)

Projection

- ❖ Removes attributes that are not in *projection list*.
- ❖ *Schema* of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- ❖ *Relational* projection operator has to eliminate *duplicates*! (Why??)
 - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (*Q*: Why not?)



sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

$\pi_{sname, rating}(S2)$

age
35.0
55.5

$\pi_{age}(S2)$

Selection

- ❖ Selects rows that satisfy a *selection condition*.
- ❖ No duplicates in result! (*Q*: Why not?)
- ❖ *Schema* of result identical to schema of its (only) input relation.
- ❖ *Result* relation can be the *input* for another relational algebra operation! (This is *operator composition*.)




sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$$\sigma_{rating > 8}(S2)$$

sname	rating
yuppy	9
rusty	10

$$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$$



Union, Intersection, Set-Difference

- ❖ All of these operations take two input relations, which must be union-compatible:
 - Same number of fields.
 - “Corresponding” fields are of the same type.
- ❖ What is the *schema* of result?

sid	sname	rating	age
22	dustin	7	45.0

$$S1 - S2$$

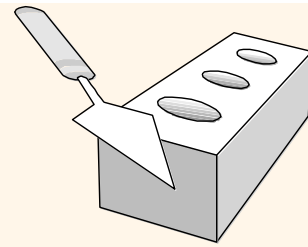
sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$$S1 \cup S2$$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$$S1 \cap S2$$

Q: Any issues w/ duplicates?



Cross-Product

- ❖ $S1 \times R1$: Each $S1$ row is paired with each $R1$ row.
- ❖ *Result schema* has one field per field of $S1$ and $R1$, with field names “inherited” if possible.
 - *Conflict*: Both $S1$ and $R1$ have a field called *sid*.

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Result
relation name

Attribute
renaming list

Source
expression E
(anything!)

- Renaming operator: $\rho (C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$



Renaming (Messy but Needed)

❖ *Conflict*: S1 and R1 both had *sid* fields, giving:

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
...
58	rusty	10	35.0	58	103	11/12/96

❖ Several renaming notations available:

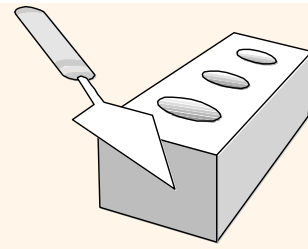
$\rho (S1R1(1 \rightarrow sid1), S1 \times R1)$ ← Positional renaming

$\rho (TempS1(sid \rightarrow sid1), S1)$ ← Name-based renaming

$TempS1 \times R1$ ← Generalized projection
(I like this one best! ☺)

$(\pi_{sid \rightarrow sid1, sname, rating, age} (S1)) \times R1$ ← SKS book's renaming

$\rho_{TempS1(sid1, sname, rating, age, sid, bid, day)} (S1 \times R1)$



Joins

❖ Condition Join: $R \bowtie_c S = \sigma_c(R \times S)$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

- ❖ *Result schema* same as that of cross-product.
- ❖ Fewer tuples than cross-product, so might be able to compute it more efficiently
- ❖ Sometimes (often!) called a *theta-join*.



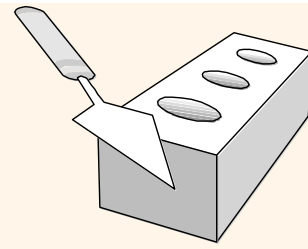
More Joins

- ❖ Equi-Join: A special case of condition join where the condition c contains only *equalities*.

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

$$S1 \bowtie_{S1.sid = R1.sid} R1$$

- ❖ Result schema is similar to cross-product, but only one copy of fields for which equality is specified.
- ❖ Natural Join: An equijoin on *all* commonly named fields (denoted simply by \bowtie).



Division

- ❖ Like join, division is not a primitive operator, but is extremely useful for expressing queries, such as:
Find sailors who have reserved all boats.
- ❖ Let A have 2 fields, x and y , while B has one field y , so we have relations $A(x,y)$ and $B(y)$:
 - **A/B contains the x tuples (e.g., sailors) such that for every y tuple (e.g., boat) in B , there is an xy tuple in A .**
 - Or: If the set of y values (boats) associated with an x value (sailor) in A contains all y values in B , the x value is in A/B .
- ❖ In general, x and y can be any lists of fields; y is the list of fields in B , and $x \cup y$ is the list of fields of A .



Examples of Division A/B

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

A

pno
p2

$B1$



sno
s1
s2
s3
s4

$A/B1$

pno
p2
p4

$B2$



sno
s1
s4

$A/B2$

pno
p1
p2
p4

$B3$




sno
s1

$A/B3$

Expressing A/B Using Basic Operators

(Advanced Topic – Just FYI ☺)



- ❖ Division not an essential op; just a useful shorthand.
(Also true of joins, but joins are so common and important that relational database systems implement joins specially.)
- ❖ *Idea*: For $A(x,y)/B(y)$, compute all x values that are not “disqualified” by some y value in B .
 - x value is *disqualified* if by attaching a y value from B , we obtain an xy tuple that *does not appear* in A .

Disqualified x values (D): $\pi_x ((\pi_x(A) \times B) - A)$

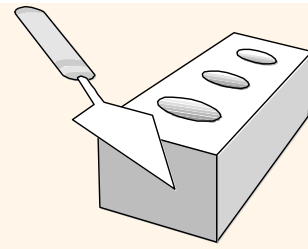
A/B : $\pi_x(A) - D$

All x 's in A

All x 's in A
combined with
all y 's in B

Tuples in the
cross product
but not in A

To Be Continued....



Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5

Reserves

sid	bid	date
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/93

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red