

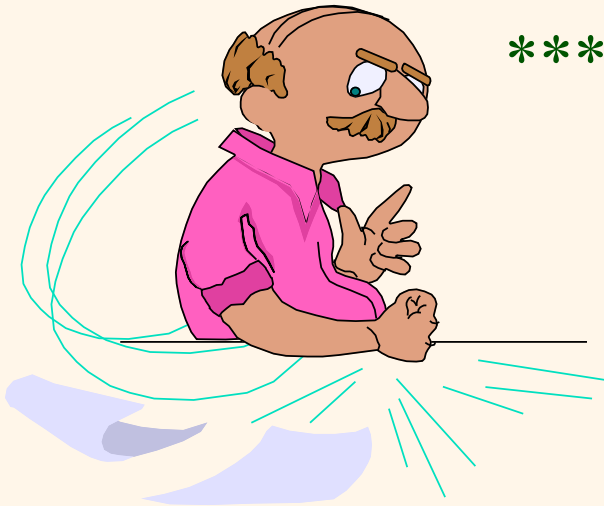


Introduction to Data Management

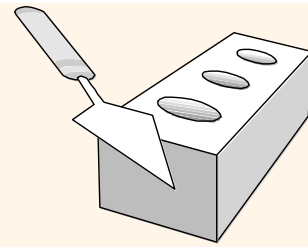
**** The “Flipped” Edition ****

Lecture #13 (SQL II)


Instructor: Mike Carey
mjcarey@ics.uci.edu



Today's Notices



- ❖ Midterm #1 is behind you!
 - Not a technical train wreck! (👍!)
 - Put it out of your minds for awhile...
- ❖ HW notes:
 - HW #2 is all graded (as you know)
 - HW #3's grading is currently in progress
 - HW #4 is your current entertainment – and we will now be on a “Friday pattern” for the next series of HW release dates and due dates
- ❖ Let's have a look at where are now, in terms of the course material...



Post-Midterm Roadmap Check

Topic Coverage and Exam Schedule


Syllabus

Topic	Reading (Required!)
Databases and DB Systems	Ch. 1
Entity-Relationship (E-R) Data Model	Ch. 6.1-6.5, 6.8-6.9
Relational Data Model	Ch. 2.1-2.4, 3.1-3.2
E-R to Relational Translation	Ch. 6.6-6.7
Relational Design Theory	Ch. 7.1-7.4.2
Midterm Exam 1	Fri, Oct 22 (during lecture time)
Relational Algebra	Ch. 2.5-2.7
Relational Calculus	⇒ Wikipedia: Tuple relational calculus
SQL Basics (SPJ and Nested Queries)	Ch. 3.3-3.5
SQL Analytics: Aggregation, Nulls, and Outer Joins	Ch. 3.6-3.9, 4.1
Advanced SQL: Constraints, Triggers, Views, and Security	Ch. 4.2, 4.4-4.5, 4.7
Midterm Exam 2	Mon, Nov 15 (during lecture time)
Storage	Ch. 12.1-12.4, 12.6-12.7
Indexing	Ch. 14.1-14.4, 14.5
Physical DB Design	Ch. 14.6-14.7, 15.1-15.3, 15.5.3
Semistructured Data Management (a.k.a. NoSQL)	Ch. 8.1, ⇒ AsterixDB SQL++ Primer , ⇒ Couchbase SQL++ Book
Data Science 1: Advanced SQL Analytics	Ch. 5.5, 11.3
Data Science 2: Notebooks, Dataframes, and Python/Pandas	Lecture notes and Jupyter notebook
Basics of Transactions	Ch. 4.3, Ch. 17
Endterm Exam	Fri, Dec 3 (during lecture time)

Midterm Exam 1

Time: Fri, Oct 22, Lecture Time
Place: SSLH 100

Nested Queries in SQL




Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid=103)
```

- ❖ A very powerful feature of SQL: a WHERE clause can itself contain an SQL query! (Actually, so can SQL's FROM and HAVING clauses!)
- ❖ To find sailors who've *not* reserved #103, use NOT IN.
- ❖ To understand semantics (including **cardinality**) of nested queries, think nested loops evaluation: *For each Sailors tuple, check qualification by computing subquery.*

Nested Queries with Correlation




Find names of sailors who've reserved boat #103:

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS (SELECT *  
               FROM Reserves R  
               WHERE R.bid=103 AND S.sid=R.sid)
```

A red curved arrow originates from the underlined 'S.sid' in the subquery's WHERE clause and points back to the 'S' in the outer query's FROM clause, illustrating the correlation between the two.

- ❖ **EXISTS** is another set comparison operator, like **IN**.
- ❖ Illustrates why, in general, a subquery must be re-computed for each Sailors tuple (conceptually).
NOTE: Recall that there was a join way to express this query, too. Relational query optimizers will try to unnest queries into joins when possible to avoid nested loop query evaluation plans.

More on Set-Comparison Operators



- ❖ We've already seen IN and EXISTS.. Can also use **NOT IN** and **NOT EXISTS**.
- ❖ Also available: *op ANY, op ALL* (for *ops*: $<$, $>$, \leq , \geq , $=$, \neq)
- ❖ Find sailors whose rating is greater than that of some sailor called Horatio:


```
SELECT *  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                      FROM Sailors S2  
                      WHERE S2.sname='Horatio')
```

So let's try ...

... running w/**ANY** on PostgreSQL

... running w/**ALL** on PostgreSQL

Rewriting *INTERSECT* Queries Using *IN*



Find sid's of sailors who've reserved both a red and a green boat:

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
      AND S.sid IN (SELECT S2.sid
                    FROM Sailors S2, Boats B2, Reserves R2
                    WHERE S2.sid=R2.sid AND R2.bid=B2.bid
                      AND B2.color='green')
```

- ❖ Similarly, *EXCEPT* queries can be re-written using *NOT IN*.
- ❖ This is what you'll have to do if using MySQL (but all the set ops are available in PostgreSQL 😊).



Division, SQL Style

Find sailors who've reserved **all** boats.

(1) SELECT S.sname
FROM Sailors **S**
WHERE NOT EXISTS

((SELECT B.bid
FROM Boats B)

EXCEPT

(SELECT R.bid
FROM Reserves R
WHERE R.sid=**S**.sid))

*(This Sailor's
unreserved
Boat ids...!)*

Sailors S such that ...

*the set of **all** Boat ids ...*

minus ...

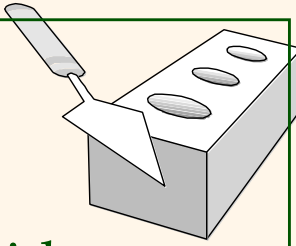
*this Sailor's
reserved Boat ids...*

*is **empty!***

Division in SQL (cont.)

(1)

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
  ((SELECT B.bid
    FROM Boats B)
  EXCEPT
  (SELECT R.bid
    FROM Reserves R
    WHERE R.sid=S.sid))
```



Find sailors who've reserved all boats.

❖ Let's do it the hard(er) way,
i.e., without EXCEPT:

(2) SELECT S.sname
FROM Sailors S

WHERE NOT EXISTS (SELECT B.bid
FROM Boats B

Sailors S such that ...


WHERE NOT EXISTS (SELECT R.bid
FROM Reserves R
WHERE R.bid=B.bid
AND R.sid=S.sid))

there is no boat B without ...

a Reserves tuple saying that S reserved B

*This way is **not** that **non-easy**
to understand – right...? (☺)*

Ordering and/or Limiting Query Results



Find the ratings, ids, names, and ages of the three best sailors

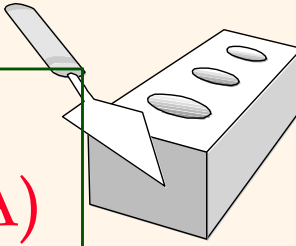
```
SELECT S.rating, S.sid, S.sname, S.age  
FROM Sailors S  
ORDER BY S.rating DESC  
LIMIT 3
```

❖ The general syntax for this:

```
SELECT [DISTINCT] expressions  
FROM tables  
[WHERE condition]  
....  
[ORDER BY expression [ ASC | DESC ]]  
LIMIT number_rows [ OFFSET offset_value ];
```

Aggregate Operators

❖ Significant extension of the relational algebra.



COUNT(*)
COUNT([DISTINCT] A)
SUM([DISTINCT] A)
AVG([DISTINCT] A)
MAX(A)
MIN(A)

single column

```
SELECT COUNT(*)  
FROM Sailors S
```


```
SELECT AVG(S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT COUNT(DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname='Bob'
```

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating= (SELECT MAX(S2.rating)  
FROM Sailors S2)
```

```
SELECT AVG(DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10 (☺?)
```

Find name and age of the oldest sailor(s)




- ❖ That first try is *illegal*!
(You'll see why shortly,
when we do **GROUP BY**.)
- ❖ *Nit*: The third version is
equivalent to the second
one, and is allowed in the
SQL/92 standard, but not
supported in some early
systems.

```
SELECT S.sname, MAX(S.age)  
FROM Sailors S
```

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE S.age =  
      (SELECT MAX(age)  
       FROM Sailors)
```

```
SELECT S.sname, S.age  
FROM Sailors S  
WHERE (SELECT MAX(S2.age)  
       FROM Sailors S2)  
      = S.age
```

Motivation for Grouping

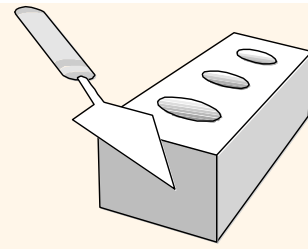


- ❖ So far, we've applied aggregate operators to **all** (qualifying) tuples. Sometimes, we want to apply them to each of several *groups* of tuples.
- ❖ Consider: *Find the age of the youngest sailor for each rating level.*
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (😊):

For $i = 1, 2, \dots, 10$:

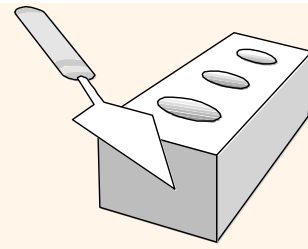
```
SELECT MIN(S.age)
FROM Sailors S
WHERE S.rating = i
```

Queries With GROUP BY and HAVING



SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

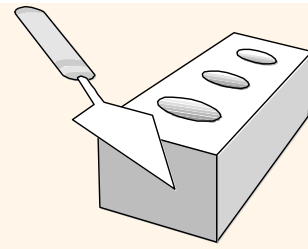
- ❖ The *target-list* contains (i) attribute names and (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
 - The attribute list (i) must be a subset of *grouping-list*.
Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a **single value per group**.
(A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)



Conceptual Evaluation

- ❖ The cross-product of *relation-list* is computed, tuples that fail the *qualification* are discarded, “unnecessary” fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- ❖ A *group-qualification* (HAVING) is then applied to eliminate some groups. Expressions in *group-qualification* must also have a *single value per group!*
 - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op must appear in *grouping-list*. (But: Some systems consider primary key semantics here.)
- ❖ One answer tuple is generated per qualifying group.

Find age of the youngest sailor with age ≥ 18
for each rating with at least 2 such sailors.



```
SELECT S.rating, MIN(S.age)
      AS minage
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT(*) >= 2
```

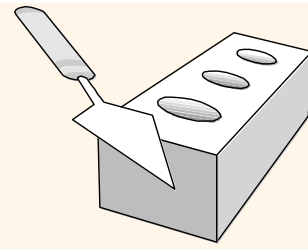
Answer relation:

rating	minage
3	25.5
7	35.0
8	25.5


Sailors instance:

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5


Find age of the youngest sailor with age ≥ 18 for each rating with at least 2 such sailors.



rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
10	16.0
9	35.0
3	25.5
3	63.5
3	25.5



rating	age
1	33.0
3	25.5
3	63.5
3	25.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0



rating	minage
3	25.5
7	35.0
8	25.5



To Be Continued...

