

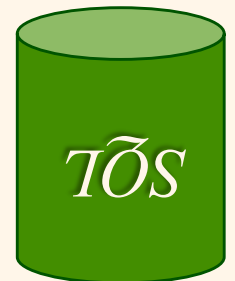



# *Introduction to Data Management*

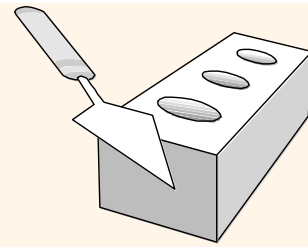
*\*\*\* The “Flipped” Edition \*\*\**

## *Lecture #16 (Advanced SQL I)*

Instructor: Mike Carey  
mjcarey@ics.uci.edu



# Announcements




- ❖ You're over half-way through...!
  - You can *do* this....! 😊
- ❖ Roadmap reminder:

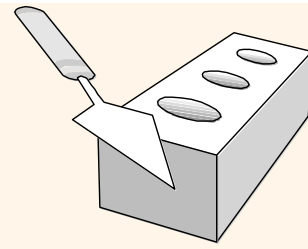
Relational Algebra	Ch. 2.5-2.7
Relational Calculus	⇒ <a href="#">Wikipedia: Tuple relational calculus</a>
SQL Basics (SPJ and Nested Queries)	Ch. 3.3-3.5
SQL Analytics: Aggregation, Nulls, and Outer Joins	Ch. 3.6-3.9, 4.1
Advanced SQL: Constraints, Triggers, Views, and Security	Ch. 4.2, 4.4-4.5, 4.7
<b>Midterm Exam 2</b>	<b>Mon, Nov 15</b> (during lecture time)

- ❖ HW#5 is in flight (and again: we're in "*Friday 6PM mode*")
  - First of the series of *SQL-based* HW assignments
  - Be sure to resolve any last PostgreSQL issues *immediately*!

# SQL Data Integrity (Largely *Review*)



- ❖ An *integrity constraint* describes a condition that every *legal instance* of a relation must satisfy.
  - Inserts/deletes/updates that violate IC's are disallowed.
  - Can be used to ensure application semantics (e.g., *sid* is a key, *bid* refers to a known boat) or prevent inconsistencies (e.g., *sname* has to be a string, integer *age* must be  $< 120$ )
- ❖ *Types of IC's*: Domain constraints, primary key constraints, foreign key constraints, unique constraints, not null constraints, general constraints.
  - *Domain constraints*: Field values must be of the right type (i.e., per the schema specification). *Always enforced!*



# SQL Data Integrity (Cont.)

❖ So far we have made good use of:

- PRIMARY KEY
- UNIQUE
- NOT NULL
- FOREIGN KEY

*Trivia Note:* MySQL will permit a “foreign key” to reference any indexed column(s)... (This enables “inclusion dependencies”.)

❖ Other features for ensuring field value integrity:

- DEFAULT (instead of NULL for missing INSERT values)
- CHECK (can access anything in the current row)


❖ More powerful integrity features include

- ~~ASSERTION~~ (book cites as unsupported, correctly 😊)
- TRIGGER (a sledge hammer to use when all else fails!)

# Some Integrity Related Examples

- ❖ CHECK is useful when more *general* ICs than just keys are involved.
- ❖ Could use SQL subqueries to express richer constraints (*if supported...* 😊).
- ❖ Constraints can be *named* (to manage them).

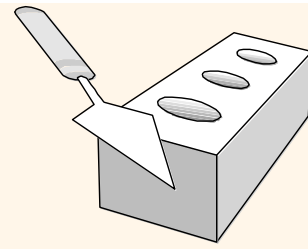
```
CREATE TABLE Sailors
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL DEFAULT 18.0,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
          AND rating <= 10 ) )
```



```
CREATE TABLE Reserves
```

```
( sid INTEGER
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid, day),
  CONSTRAINT noInterlakeRes
  CHECK ( 'Interlake' <>
    (-SELECT B.bname
    FROM Boats B
    WHERE B.bid=bid)))
```






# *Quick PostgreSQL Examples*

```
ALTER TABLE Sailors  
ALTER COLUMN age SET DEFAULT 18.0;
```


```
ALTER TABLE Sailors  
ADD CONSTRAINT ValidRatingConstraint  
CHECK (rating >= 1 AND rating <= 10);
```

```
ALTER TABLE Sailors  
ADD CONSTRAINT AliveConstraint  
CHECK (age > 0);
```

```
INSERT INTO SAILORS (sid, sname, rating, age)  
VALUES (110, 'Mike', 9, 0.0);
```




# Enforcing Referential Integrity (RI)



- ❖ Consider Sailors and Reserves; *sid* in Reserves is a foreign key (FK) that references Sailors.
- ❖ What should be done if a Reserves tuple with a non-existent sailor id is *inserted*? (**A: Reject it!**)
- ❖ What should be done if a Sailors tuple is *deleted*?
  - Also delete all Reserves tuples that refer to it, or
  - Disallow deletion of a Sailor that's being referred to, or
  - Set *sid* in Reserves tuples that refer to it to some *default sid*.
  - Could opt to set *sid* in Reserves tuples that refer to it to *null*, but this would not be a great idea (or allowed if key).
- ❖ Similar issues if a Sailor's primary key is *updated*!





# RI Enforcement in SQL (Reminder)

- ❖ SQL/92 and SQL:1999 support all 4 options on deletes and updates.
  - Default is **NO ACTION** (delete/update is rejected)
  - **CASCADE** (also delete all tuples that refer to the deleted tuple)
  - **SET NULL / SET DEFAULT** (set foreign key value of referencing tuple)

Ex:

```
CREATE TABLE Reserves  
(sid INTEGER,  
bid INTEGER,  
date DATE,
```


....

```
FOREIGN KEY (sid)  
REFERENCES Sailors  
ON DELETE CASCADE  
ON UPDATE SET NULL)
```

***Note:** An odd combo; just illustrating some of what's possible...*



# *Stored Procedures (and Functions) in SQL*

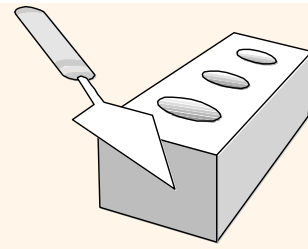


## ❖ What is a stored procedure?

- A program executed via a single SQL statement
- Executes in the process space of the **server**

## ❖ Advantages:

- Can encapsulate application logic while staying “close” to the data
- Supports the reuse (sharing) of application logic by different users
- Can be used to help secure database applications, as we will see a bit later on



# Stored Procedures: More Detail

- ❖ A **stored procedure** is a function or procedure written in a *general-purpose* programming language that executes *within* the DBMS.
- ❖ These can perform computations that *cannot* be expressed in SQL – i.e., they go *beyond* the limits of relational completeness.
- ❖ Procedure execution is requested through a single SQL statement (**CALL**).
- ❖ Executes on the (usually *remote*!) DBMS server.
- ❖ SQL **PSM** (*Persistent Stored Modules*) extends SQL with concepts from general-purpose PLs.



# Stored Procedures: *Functions*

Ex: Let's define a simple **function** that we might want:

```
CREATE FUNCTION ShowNumReservations(boat_id int)
RETURNS TABLE (sid int, sname text, cnt bigint)
AS $$
```

```
    SELECT S.sid, S.sname, COUNT(*) AS cnt
    FROM Sailors S, Reserves R
    WHERE S.sid = R.sid AND R.bid = boat_id
    GROUP BY S.sid;
```

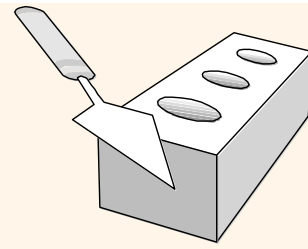
```
$$
```

```
LANGUAGE SQL;
```

**Q:** *What does this function do?*

*First:* INSERT INTO Reserves VALUES – need more data  
(22, 101, '2021-10-31'), (22, 104, '2021-10-31');

*Then:* SELECT \* FROM ShowNumReservations (101);



# Stored Procedures: *Procedures*

Ex: Let's define a **procedure** that might be useful:

(Possible modes for parameters: IN, OUT, INOUT)

```
CREATE PROCEDURE IncreaseRating (  
    IN sailor_sid int, IN increase int)
```

```
AS $$
```

```
    UPDATE Sailors
```

```
    SET rating = rating + increase
```

```
    WHERE sid = sailor_sid;
```


```
$$
```

```
LANGUAGE SQL;
```

**Q:** *How is this **procedure** different?*

*Then:* **CALL IncreaseRating(95,1);**


# Stored Procedures: *External Logic*



Stored procedures can also be written outside of the SQL language:

```
CREATE PROCEDURE RecklessSailors ( )  
AS 'DIRECTORYfuncs', 'reckless_sailors'  
LANGUAGE C;
```


# Main *SQL/PSM* Constructs (FYI)



- ❖ Supports FUNCTIONS and PROCEDUREs
- ❖ Local variables (**DECLARE**)
- ❖ RETURN values for FUNCTION
- ❖ Assign variables with **SET**
- ❖ Branches and loops:
  - IF (condition) THEN statements;  
ELSEIF (condition) statements;  
... ELSE statements; END IF;
  - LOOP statements; END LOOP
- ❖ Queries can be parts of expressions
- ❖ Cursors available to iterate over query results

**Note:** *SQL PSM is the SQL standard language for S.P.'s; **not** supported by all vendors – and **not** in PostgreSQL – due to the relative **lateness** of its standardization...!)*

# A (random 😊) SQL/PSM Example




```
CREATE FUNCTION ResvRateSailor (IN sailorId INT)
  RETURNS INT
BEGIN
  DECLARE resvRating INT
  DECLARE numResv INT
  SET numResv = (SELECT COUNT(*)
                 FROM Reserves R
                 WHERE R.sid = sailorId)
  IF (numResv > 10) THEN resvRating = 1;
    ELSE resvRating = 0;
  END IF;
  RETURN resvRating;
END;
```

**Note:** See your chosen RDBMS's docs for info about its procedural extension to SQL...




# Triggers in SQL



- ❖ Trigger: a procedure that runs automatically if specified changes occur to the DBMS
- ❖ Three parts:
  - Event (activates the trigger)
  - Condition (tests if the trigger should run)
  - Action (what happens if the trigger runs)
- ❖ Can be used to do “whatever”!
  - In PostgreSQL, “whatever” is a stored function call; it can also cause the current update to bail out.
  - Details vary *WIDELY* from vendor to vendor (!)
  - Major source of “vendor lock-in”, along with their *stored procedure language* (= trigger action language)

# Trigger Syntax (*PostgreSQL*)



## CREATE TRIGGER

CREATE TRIGGER — define a new trigger

## Synopsis

```
CREATE [ CONSTRAINT ] TRIGGER name { BEFORE | AFTER | INSTEAD OF } { event [ OR ... ] }  
ON table_name  
[ FROM referenced_table_name ]  
[ NOT DEFERRABLE | [ DEFERRABLE ] [ INITIALLY IMMEDIATE | INITIALLY DEFERRED ] ]  
[ REFERENCING { { OLD | NEW } TABLE [ AS ] transition_relation_name } [ ... ] ]  
[ FOR [ EACH ] { ROW | STATEMENT } ]  
[ WHEN ( condition ) ]  
EXECUTE { FUNCTION | PROCEDURE } function_name ( arguments )
```

where **event** can be one of:

```
INSERT  
UPDATE [ OF column_name [, ... ] ]  
DELETE  
TRUNCATE
```

```
CREATE TRIGGER trigger_name  
{BEFORE | AFTER} { event }  
ON table_name  
[FOR [EACH] { ROW | STATEMENT }]  
EXECUTE PROCEDURE trigger_function
```



(See <https://www.postgresql.org/docs/14/sql-createtrigger.html>)

# Trigger Example (*PostgreSQL*)



-- First create the logic, i.e., the “whatever” part

```
CREATE FUNCTION AddYoungSailor ( )  
RETURNS Trigger  
AS $$  
    BEGIN  
        INSERT INTO YoungSailors (sid, sname, age, rating)  
            VALUES (NEW.sid, NEW.sname, NEW.age, NEW.rating);  
        RETURN NEW;  
    END;  
$$  
LANGUAGE PLPGSQL;
```



# Trigger Example (*PostgreSQL, cont.*)




-- Now create the trigger itself

```
CREATE TRIGGER YoungSailorLogger  
AFTER INSERT ON Sailors  
FOR EACH ROW ←  
WHEN (NEW.age < 18)  
EXECUTE FUNCTION AddYoungSailor ( );
```

Note: *FOR EACH ROW* provides less power than *FOR EACH STATEMENT* (e.g., can't compute average new age)

# Trigger Example (PostgreSQL, cont.)



**FIRST:** CREATE TABLE YoungSailors (LIKE Sailors);

And now let's try some INSERTs:

- ☐ INSERT INTO Sailors(sid, sname, rating, age)  
VALUES (777, 'Lucky', 7, 77);
- ☒ INSERT INTO Sailors(sid, sname, rating, age)  
VALUES (778, 'Lucky Jr', 7, 17);

*(NOTE: Look at **YoungSailors** table content after each one!)*

# Trigger Example (PostgreSQL, cont.)



-- Let's implement a poor man's CHECK constraint!

```
CREATE FUNCTION BlockYoungSailor ( ) RETURNS Trigger AS $$
```

```
BEGIN
```

```
    RAISE 'Sailors must be at least 10';
```

```
END;
```

```
$$ LANGUAGE PLPGSQL;
```

```
CREATE TRIGGER SailorAgeEnforcer
```

```
BEFORE INSERT ON Sailors
```

```
FOR EACH ROW
```

```
WHEN (NEW.age < 10)
```

```
EXECUTE FUNCTION BlockYoungSailor ( );
```

```
INSERT INTO Sailors (sid, sname, rating, age)  
VALUES (800, 'Baby Face', 10, 1);
```

*To Be Continued...*

