





Introduction to Data Management

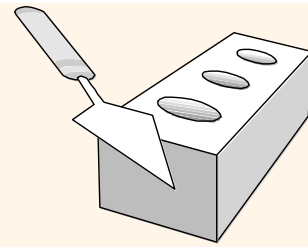
**** The “Flipped” Edition ****

Lecture #18 *(Storage & Indexing I)*

Instructor: Mike Carey
mjcarey@ics.uci.edu



Announcements



❖ Roadmap check:


Midterm Exam 2 ←	Mon, Nov 15 (during lecture time)
Storage ←	Ch. 12.1-12.4, 12.6-12.7
Indexing	Ch. 14.1-14.4, 14.5
Physical DB Design	Ch. 14.6-14.7, 15.1-15.3, 15.5.3
Semistructured Data Management (a.k.a. NoSQL)	Ch. 8.1, → AsterixDB SQL++ Primer , → Couchbase SQL++ Book
Data Science 1: Advanced SQL Analytics	Ch. 5.5, 11.3
Data Science 2: Notebooks, Dataframes, and Python/Pandas	Lecture notes and Jupyter notebook
Basics of Transactions	Ch. 4.3, Ch. 17
Endterm Exam	Fri, Dec 3 (during lecture time)

- ❖ SQL HW assignments (in “Friday 6PM mode”)
 - HW #5 (Basic SQL) is due on Friday
 - HW #6 (Advanced SQL) will come out on Friday
- ❖ Midterm #2 (relational languages) is *1.5 weeks away (!)*
 - Same in-class Gradescope + cheat sheet logistics as Midterm #1

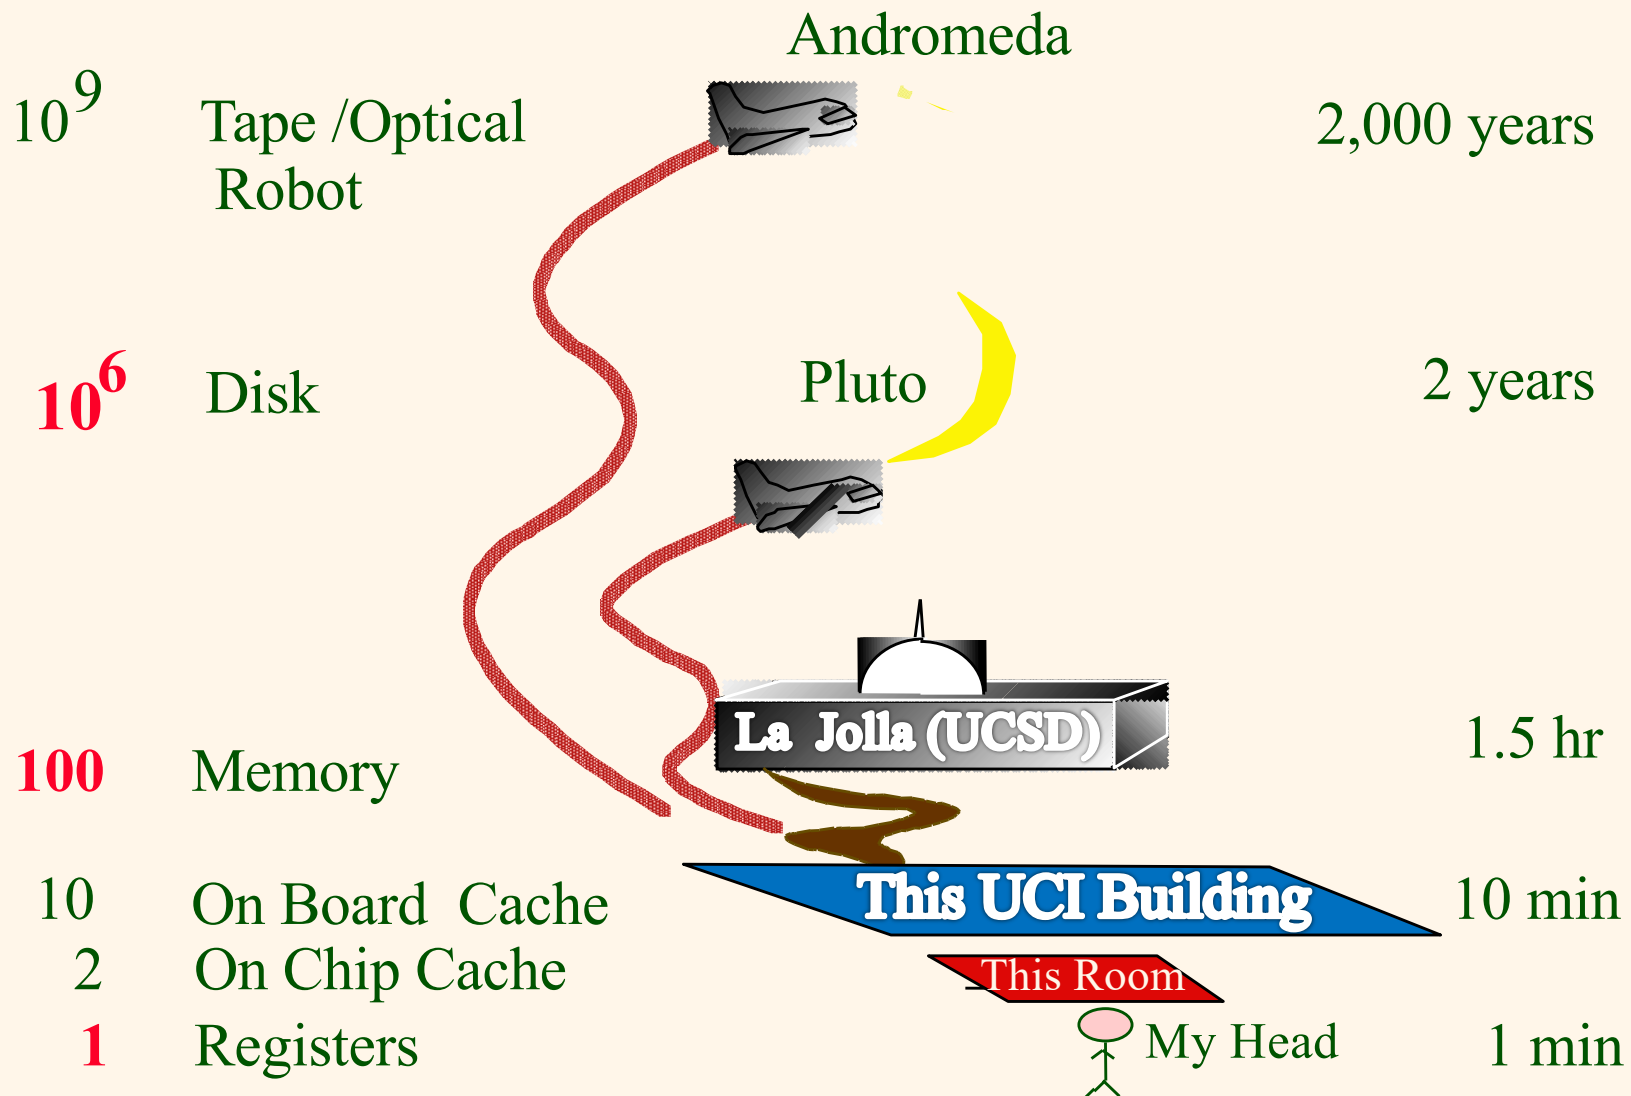



Disks and Files

- ❖ DBMSs store data on *secondary storage*.
- ❖ This has major implications for DBMS design!
 - **READ:** xfer data from *disk* to *memory* (RAM).
 - **WRITE:** xfer data from RAM to disk (HDD or SSD).
 - Both are high-cost operations, relative to in-memory operations, so must be considered carefully!



Storage Hierarchy & Latency (Jim Gray in the mid 1990's): How Far Away is my Data?






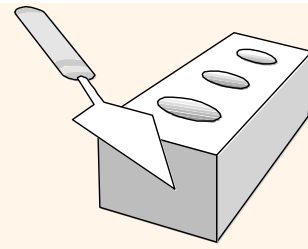
Why Not Store Data in Main Memory?

- ❖ *Main memory (RAM) costs too much! Roughly:*
 - RAM: \$22/GB [Dell @ 05/2019]
 - SSD: \$1/GB (22x cheaper than RAM)
 - Disk: \$0.16/GB (138x cheaper than RAM, 5.5x vs. SSD)
- ❖ *Main memory is **volatile**.* We want our data to be saved between runs. (Obviously...!)
- ❖ Your typical (basic) storage hierarchy:
 - Main memory (RAM) for currently used data
 - Disk (HDD, or increasingly SSD) for the main database (secondary storage)
 - Tapes for archiving the data (tertiary storage)

Components of a Hard Disk


- ❖ The platters spin (10,000 rpm)
- ❖ The arm assembly is moved in or out to position a head on a desired track
Tracks under heads form a *cylinder* (imaginary!)
- ❖ Only one head reads/writes at any one time.
- ❖ *Block size* is a multiple of *sector size* (which is fixed)





Accessing a Disk Page

- ❖ Time to access (read/write) a disk block:
 - *Seek time* (moving arms to position disk head on track)
 - *Rotational delay* (waiting for block to rotate under head)
 - *Transfer time* (actually moving data to/from disk surface)
- ❖ Seek time and rotational delay dominate!
 - Seek time varies from about 1 to 20msec
 - Rotational delay varies from 0 to 10msec
 - Transfer rate is < 1 msec per 4KB page
 - Key to I/O cost: **Reduce seek/rotation delays!**
 - ➔ **Bottom line: *Random vs. sequential I/O***




What About SSD Storage?


- ❖ NAND Flash: *block-oriented* interface like HDD
 - *Random reads*: 20-100 usec (vs. 5-10 msec for HDD)
 - *Transfer rate*: 500-3000 GB/sec (vs. 200 for HDD)
 - *Read performance*: 10,000+ IOPS (random reads/sec), but up to 10x with parallel accesses (unlike HDD)
 - ❖ Writes are more complicated than HDD
 - *Random writes*: ~100 usec (first time in empty spot)
 - *Write performance*: 10,000+ IOPS (random reads/sec)
 - Limited life, wear leveling, flash translation layer, ...
- Bottom line: No seeks, but still block-oriented**



So: Disks and Files (Reminder!)

- ❖ DBMSs store persistent data on *secondary storage*.
- ❖ This has major implications for DBMS design!
 - **READ:** xfer data block from *disk* to *memory* (RAM).
 - **WRITE:** xfer data block from RAM to disk.
 - Both are *high-cost* operations, relative to in-memory operations, so must be considered carefully!





Ex: Emp(eid, ename, sal, deptid)

P1

1	111	Smith	3K	1
2	222	Lee	100K	3
3	333	Carey	80K	1
4	444	Smith	12K	7

P2

1	555	Smith	18K	3
2	666	Jones	90K	5
3	777	Smith	23K	4
4	888	Krishan	60K	8

← Record id
(RID) is (P2,3)

Underlying
Emp
file pages

P10000

...

1
2
3
4	9999999	Smith	18K	11



Processing a Query

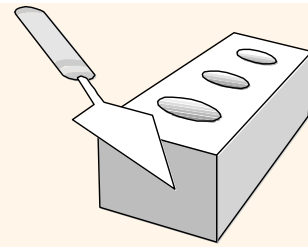
❖ Suppose someone asks a simple SQL query:

```
SELECT * FROM Emp WHERE eid = 12345;
```

❖ Some processing options would include:


- *Option 1:* Sequentially scan the data file (and stop, if we know eid is a key) → 5000 page reads (*avg.*)
- *Option 2:* Binary search the data file (and stop, if we know eid is a key) → $\log_2(10,000) \approx 15$ page reads (*avg.*)
- Even though Option 2 is $\approx 300\times$ faster, we'd like to do *even better* (especially for large data sets!)

Indexing is the Answer!



Key value k or key
range (k_1, k_2)

777



Index on k



$I(k)$ or
 $I(k_1), \dots, I(k_2)$

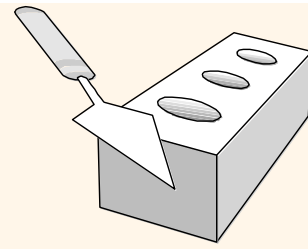
(P2,3)

- ❖ *Index* maps from *keys* to associated *info I*
 - $I(k)$ can be the *data record itself* with key k , or
 - $I(k)$ can be the *RID* of the data record with key k , or
 - $I(k)$ can be a *list of RIDs* of data records with key k !
 - Alternatively, we could map from data field values to the *PK value(s)* of the associated record(s) – SQL Server and AsterixDB do this, for example.





Indexes

- ❖ An *index* on a file speeds up selections on the *search key fields* for the index.
 - Any subset of the fields of a relation can serve as the search key for an index on the relation.
 - *Search key* is **not** the same as a “*key*” (i.e., it’s not the primary key, it’s just a field we’re very interested in).
- ❖ An index contains a collection of *data entries*, and it supports efficient retrieval of *all* data entries k^* with a given key value k .
 - Given a **data** entry k^* , we can find **1st actual record** with key k with ~ 1 more disk I/O. (Details soon...)

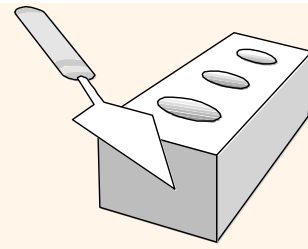


Ex: *Emp* (*eid*, *ename*, *sal*, *deptid*)

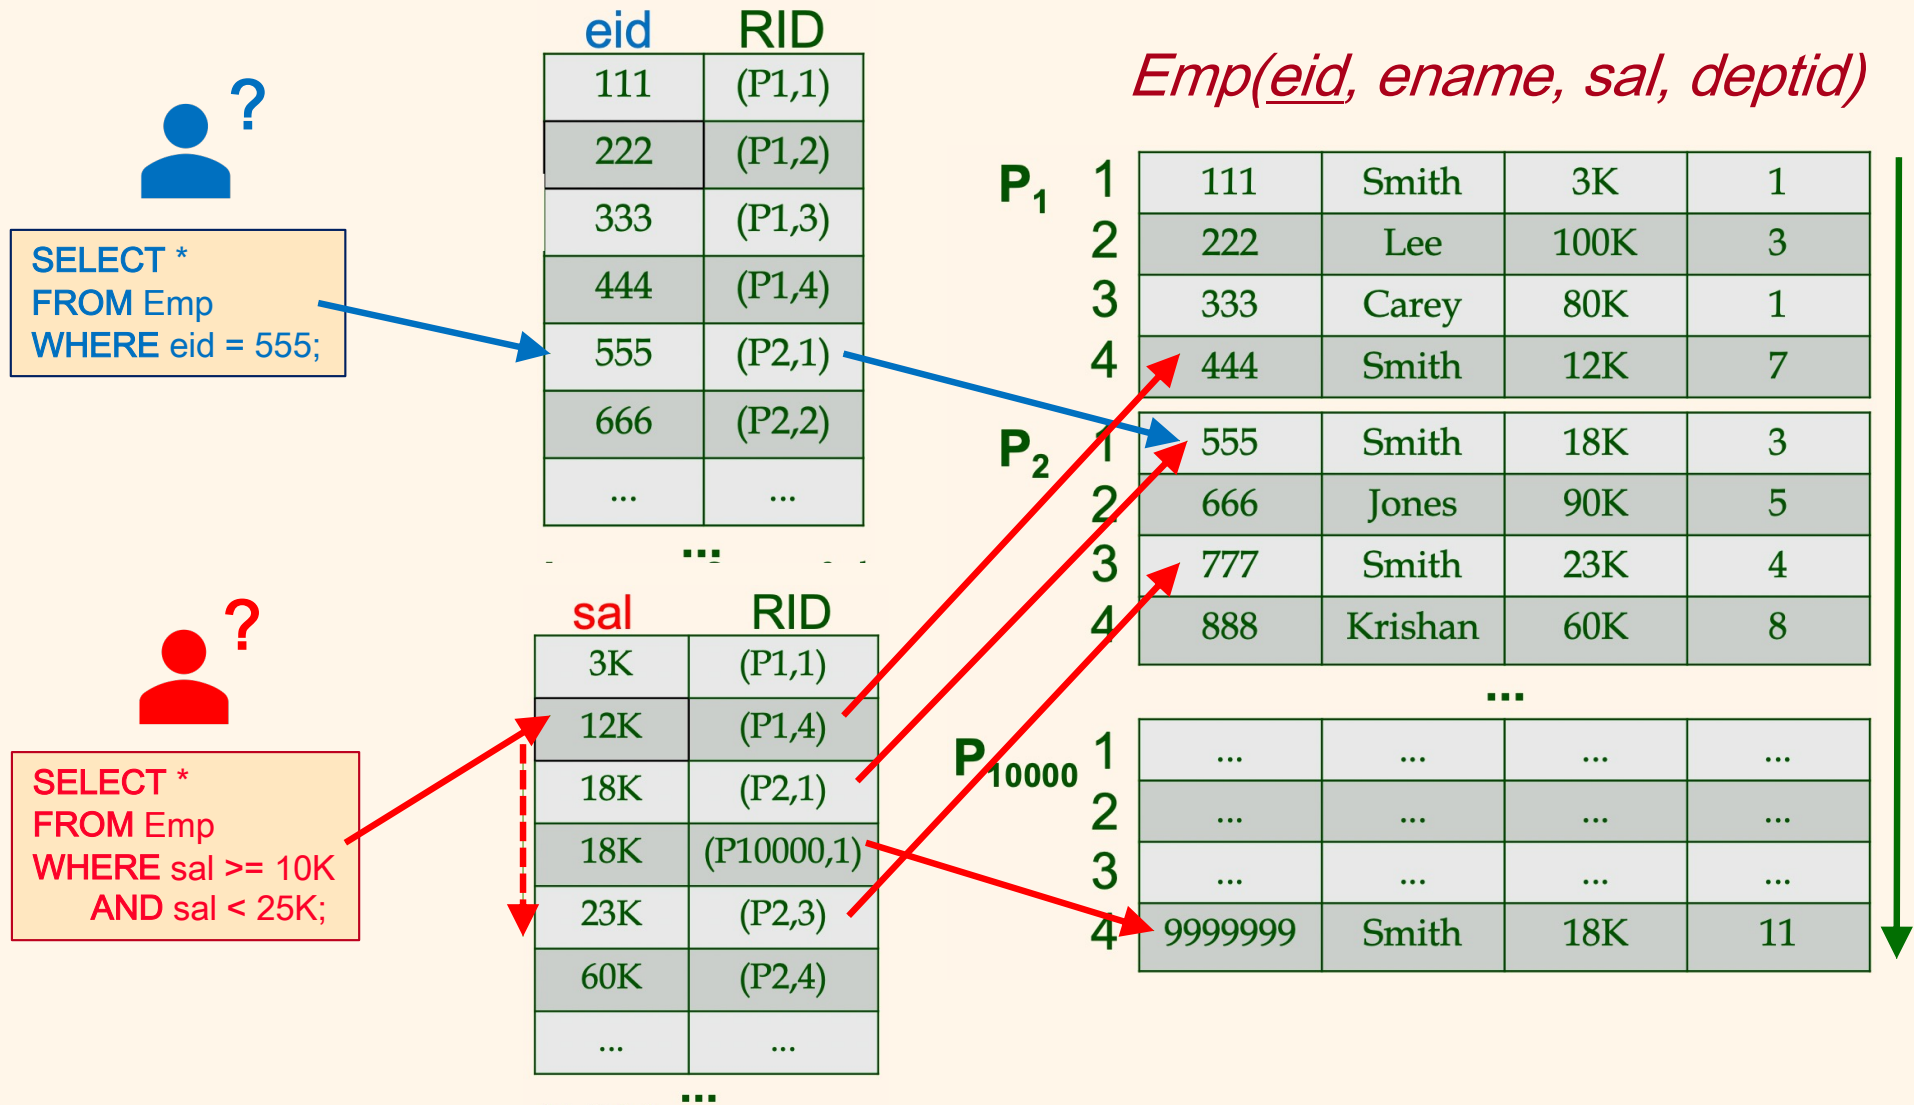
- ❖ One simple approach would be to have one more file, sorted *on k*, for each *k* that we want to index
 - *Hundreds* of (key, RID) entries will fit on a *single page*
 - Index is thus *much* smaller than the data file
 - Less data (*fewer reads!*) to search to locate the RIDs of interest

<i>eid</i>	RID		<i>sal</i>	RID	
111	(P1,1)	 (RID)	3K	(P1,1)	 (RID)
222	(P1,2)		12K	(P1,4)	
333	(P1,3)		18K	(P2,1)	
444	(P1,4)		18K	(P10000,1)	
555	(P2,1)		23K	(P2,3)	
666	(P2,2)		60K	(P2,4)	
...	


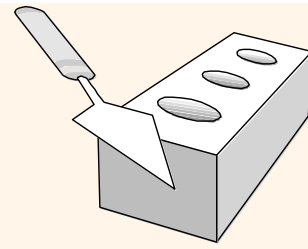
Note:
*Can have
multiple
unclustered
indexes!*



Indexed Tables Under the Hood




Even Better: Tree Indexes!

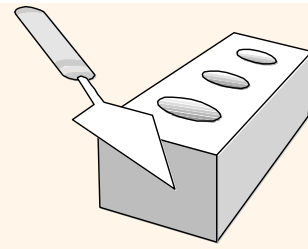


- ❖ Leaf pages contain *data entries*, and are chained
- ❖ Non-leaf pages have *index entries*; role is to guide searches
- ❖ Query processing steps become:
 1. Choose a good index to use (if one is available)
 2. Search the index to determine the interesting RID(s)
 3. Use the RID(s) to fetch the corresponding record(s)

Tree-Structured Indexes: Over(re)view

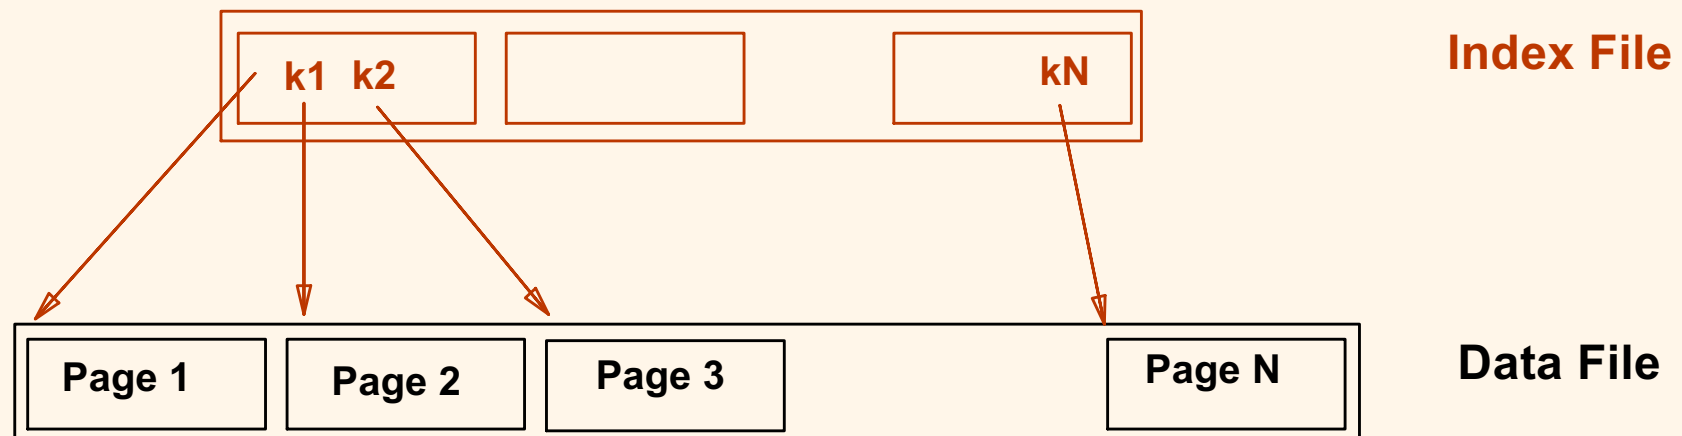


- ❖ As for *any* index, 3 alternatives for data entries \mathbf{k}^* :
 - Data record with key value \mathbf{k}
 - $\langle \mathbf{k}, \text{RID of data record with search key value } \mathbf{k} \rangle$
 - $\langle \mathbf{k}, \text{list of RIDs of data records with search key } \mathbf{k} \rangle$
- ❖ This data entry choice is orthogonal to the *indexing technique* used to locate data entries \mathbf{k}^* .
- ❖ Tree-structured indexing techniques support both *range searches* and *equality searches*.
- ❖ ISAM: static structure; B+ tree: dynamic, adjusts gracefully under inserts and deletes.





Range Searches (Review)

- ❖ “Find all students with $\text{gpa} > 3.0$ ”
- ❖ If records are in a sorted file, do binary search to find first such student, then scan to find others.
 - Cost of file binary search can be quite high.
- ❖ Simple idea to do better: add an “index” file.

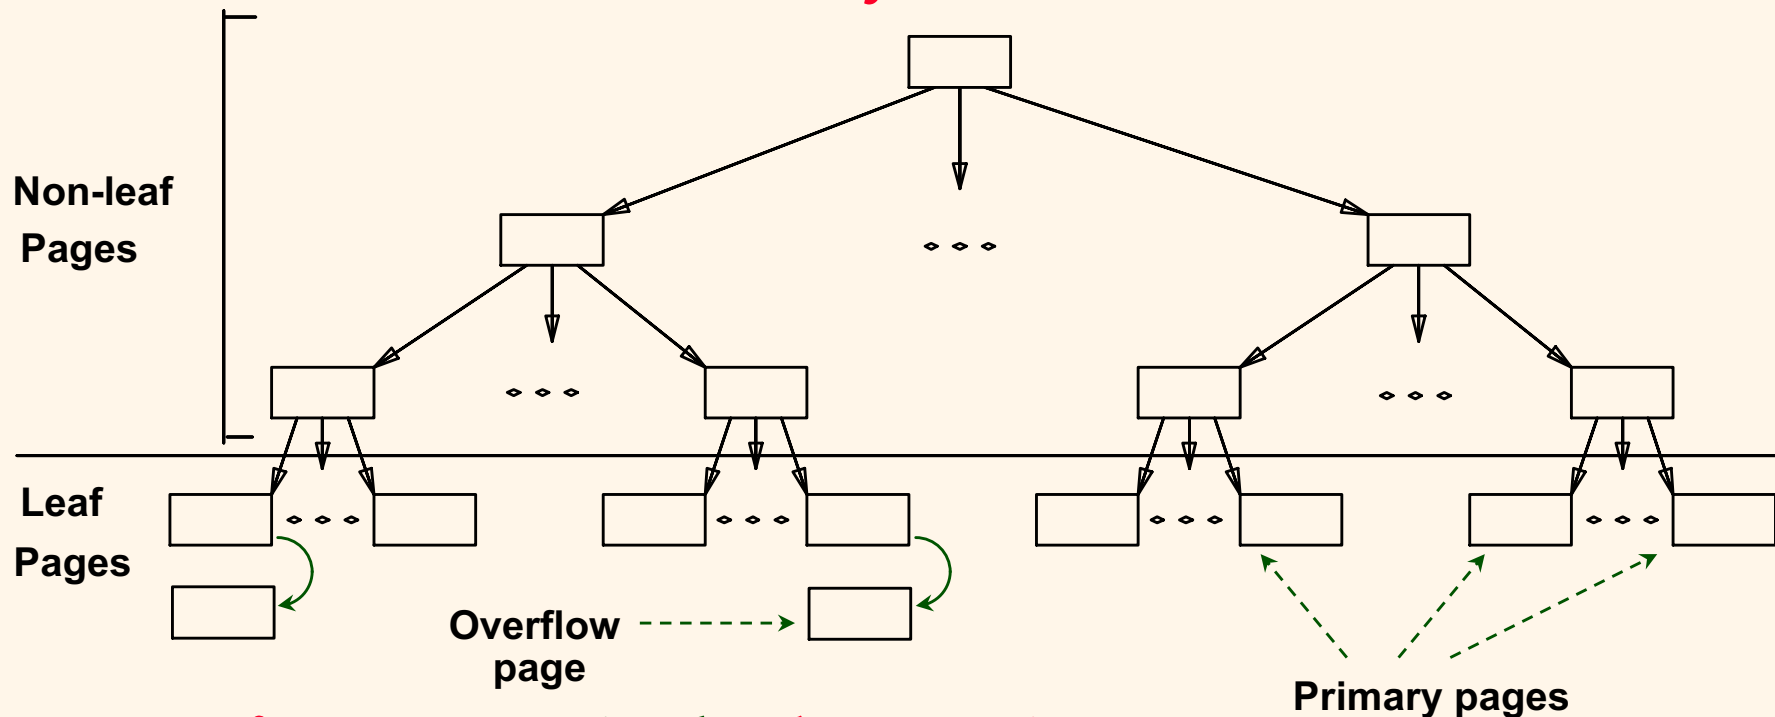


➡ *Can now binary search the (smaller) index file!*

ISAM* index entry



- ❖ Index file may still be quite large. But, we can apply the same idea **recursively** to address that!




👉 *Leaf pages contain the data entries.*

**ISAM: Indexed Sequential Access Method*

ISAM Operation(s)


- ❖ *File creation*: Leaf (data) pages first allocated sequentially, sorted by search key; then index pages allocated, and then overflow pages.
- ❖ *Index entries*: $\langle \text{search key value, page id} \rangle$; they “direct” searches for *data entries*, which are in leaf pages.
- ❖ *Search*: Start at root; use key comparisons to go to leaf.
 $\text{I/O cost} \propto \log_F N$; $F = \# \text{entries/index pg}$, $N = \# \text{leaf pgs}$
- ❖ *Insert*: Find leaf data entry belongs to, put entry there.
- ❖ *Delete*: Find and remove entry from leaf; if overflow page is empty now, de-allocate it.
- ➡ **Static tree structure**: *inserts/deletes affect only leaf pages.*



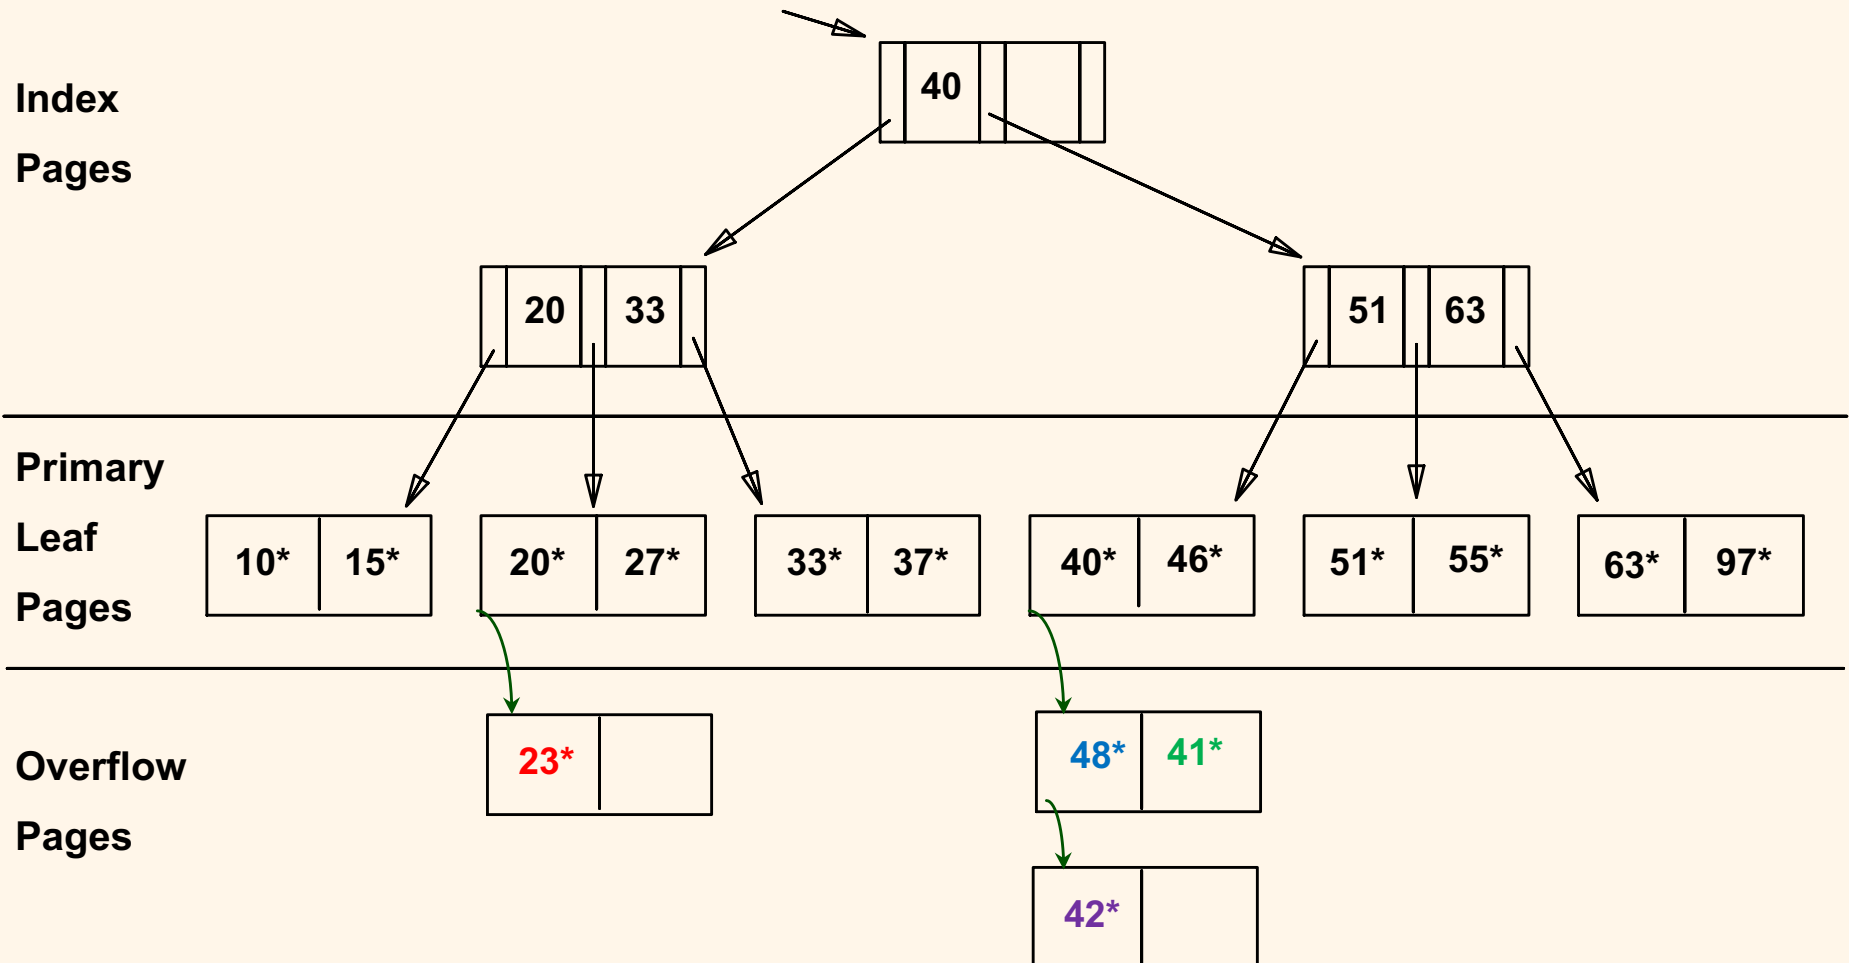



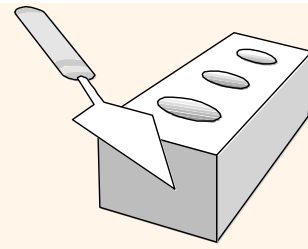
Example ISAM Tree

- ❖ Suppose each node can hold 2 entries (really more like 200, since nodes are disk pages!)




After Inserting 23, 48*, 41*, 42* ...*





... Then Deleting 42*, 51*, 97*



➡ Note that 51* still appears in index levels, but **not** in leaf!

To Be Continued...

