



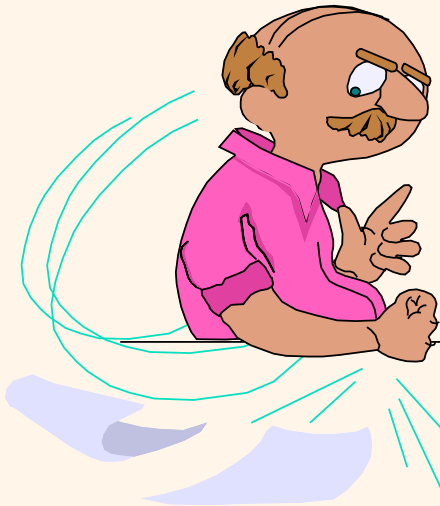
# *Introduction to Data Management*

*\*\*\* The “Flipped” Edition \*\*\**

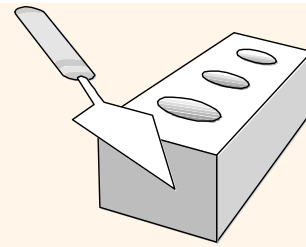
## *Lecture #20*

*(Storage & Indexing III  
and Physical DB Design I)*


Instructor: Mike Carey  
mjcarey@ics.uci.edu



# Announcements



## ❖ Roadmap reminder:

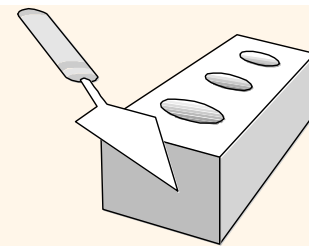
<b>Midterm Exam 2</b> 	<b>Mon, Nov 15</b> (during lecture time)
Storage	Ch. 12.1-12.4, 12.6-12.7
Indexing	Ch. 14.1-14.4, 14.5
Physical DB Design	Ch. 14.6-14.7, 15.1-15.3, 15.5.3
Semistructured Data Management ( <i>a.k.a.</i> NoSQL)	Ch. 8.1, <a href="#">⇨ AsterixDB SQL++ Primer</a> , <a href="#">⇨ Couchbase SQL++ Book</a>
Data Science 1: Advanced SQL Analytics	Ch. 5.5, 11.3
Data Science 2: Notebooks, Dataframes, and Python/Pandas	Lecture notes and Jupyter notebook
Basics of Transactions	Ch. 4.3, Ch. 17
<b>Endterm Exam</b>	<b>Fri, Dec 3</b> (during lecture time)

- ❖ HW #6 should be wrapping up now!
  - Second in the series of *SQL-based* HW assignments
  - Due **this Friday @ 6 PM** (w/ usual 24-hour late window)
- ❖ Midterm #2 is now one weekend away (*wow!*)
  - Monday (Nov. 15), conducted just like Midterm #1
  - In person, Gradescope + hard copy cheat sheet, assigned seats
- ❖ **Today:** Finish Storage & Indexing, then move to Physical DB Design!



# *A Note on B+ Tree “Order”*

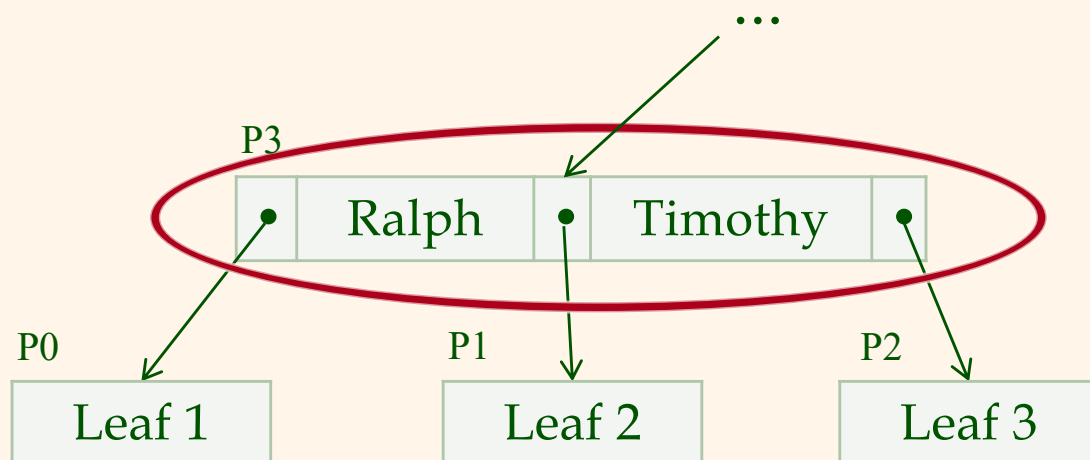
- ❖ (Mythical!) *order (d)* concept replaced by physical space criterion in practice (“*at least half-full*”).
  - Index pages can typically hold many more entries than leaf pages.
  - Variable-sized records and search keys mean that different nodes will contain different numbers of entries.
  - Even with fixed length fields, multiple records with the same search key value (*duplicates*) can lead to variable-sized data entries in the tree’s leaf pages.



# (Page Implementation *Details*)

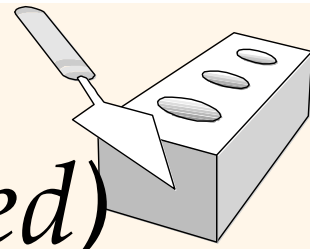
**Q:** What if you were to “open up” a *B+ Tree page*?

- Control info (e.g., level, # children, free space offset)
- Search key array (with possible on-page indirection for variable-length data, using offsets), or a key/data array – for non-leaf *vs.* leaf pages, respectively
- Child pointer array, where pointer = page id on disk!



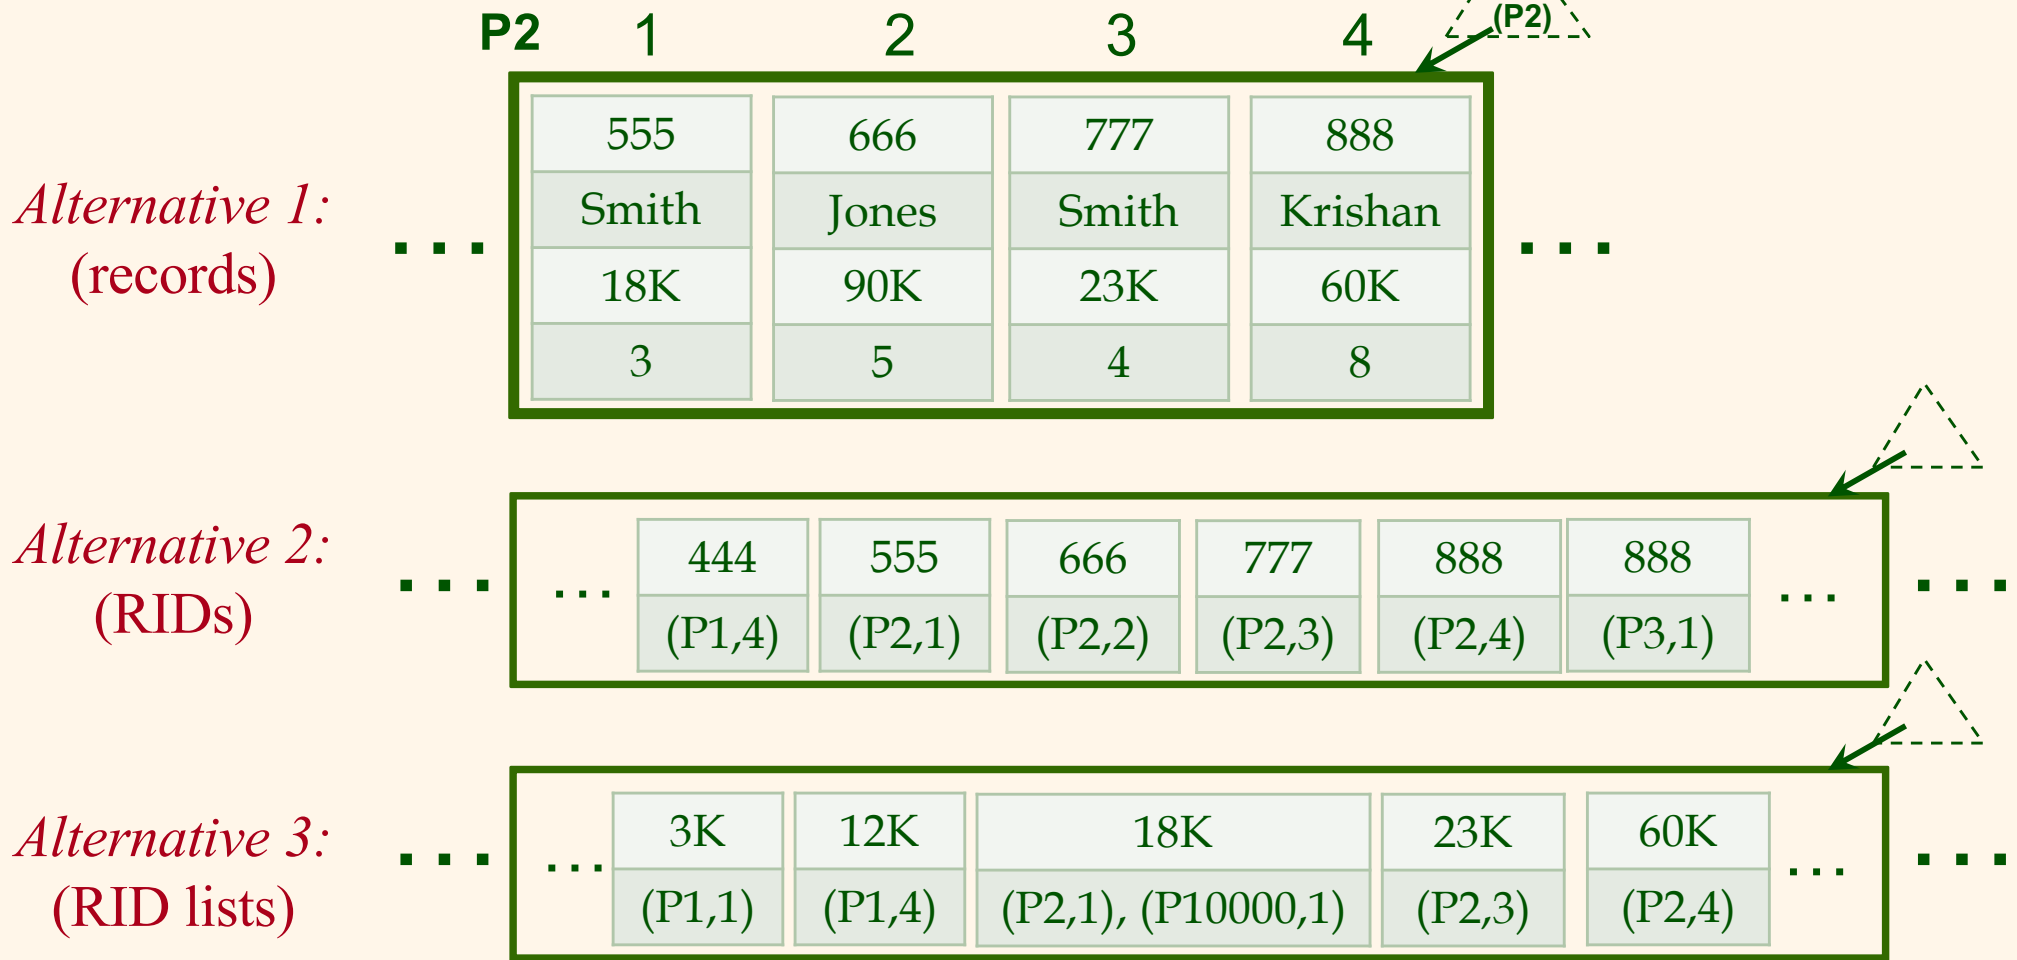
0	Level (1)
	NumChildren (3)
8	Free offset (40)
	Key 0 offset (32)
	Key 1 offset (37)
20	Child 1 page id (P0)
	Child 2 page id (P1)
	Child 3 page id (P2)
32	Key 0 ("Ralph")
37	Key 1 ("Timothy")
40	...

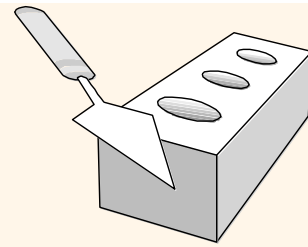
} *not to scale...*



# (Leaf Page I(k) Alternatives Revisited)

Ex: Emp(eid, ename, sal, deptid)





# (Leaf Page I(k) Alternatives, cont.)

Ex: Emp(eid, ename, sal, deptid)

**P2**      1                      2                      3                      4                      (P2)

555	666	777	888
Smith	Jones	Smith	Krishnan
18K	90K	23K	60K
3	5	4	8

Alternative 1:  
(records)

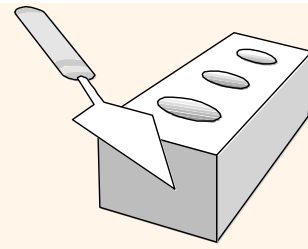
Note: Must use PKs  
in secondary indexes  
if the primary index  
uses Alternative 1!  
(Think about why!)

Alternative 2':  
(PKs)

...	444	555	666	777	888	888	...
...	444	555	666	777	888	999	...

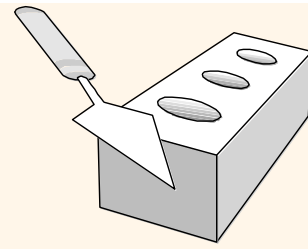
Alternative 3':  
(PK lists)

...	3K	12K	18K	23K	60K	...
...	111	444	555, 4439667	777	888	...



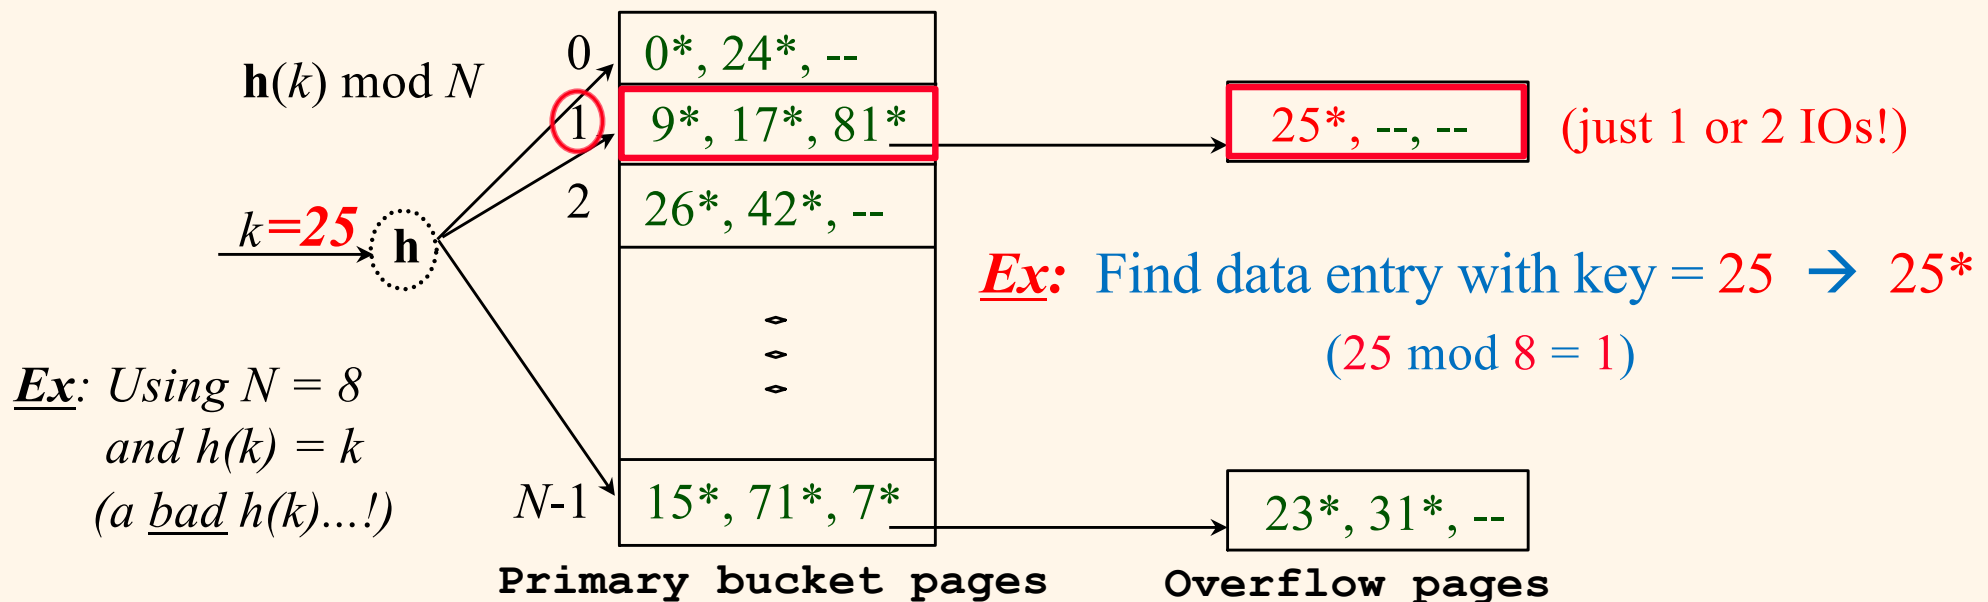
# *Hash-Based Indexes*

- ❖ Hash-based indexes are fast for *equality selections*.  
*Cannot* support range searches.
- ❖ Static and dynamic hashing techniques exist; trade-offs similar to *ISAM* vs. B+ trees.
- ❖ *As for any index, 3 alternatives for data entries  $k^*$ :*
  - Data record with key value  $k$
  - $\langle k, \text{rid of data record with search key value } k \rangle$
  - $\langle k, \text{list of rids of data records with search key } k \rangle$
  - Choice is orthogonal to the *indexing technique*!



# Static Hashed Indexes

- ❖ # primary pages fixed, allocated sequentially, never de-allocated; overflow pages if needed.
- ❖  $h(k) \bmod N = \text{bucket (page) to which data entry with key } k \text{ belongs. (} N = \# \text{ of buckets)}$

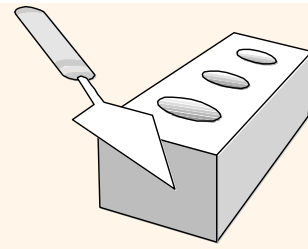






## *Static Hashed Indexes (Cont'd.)*

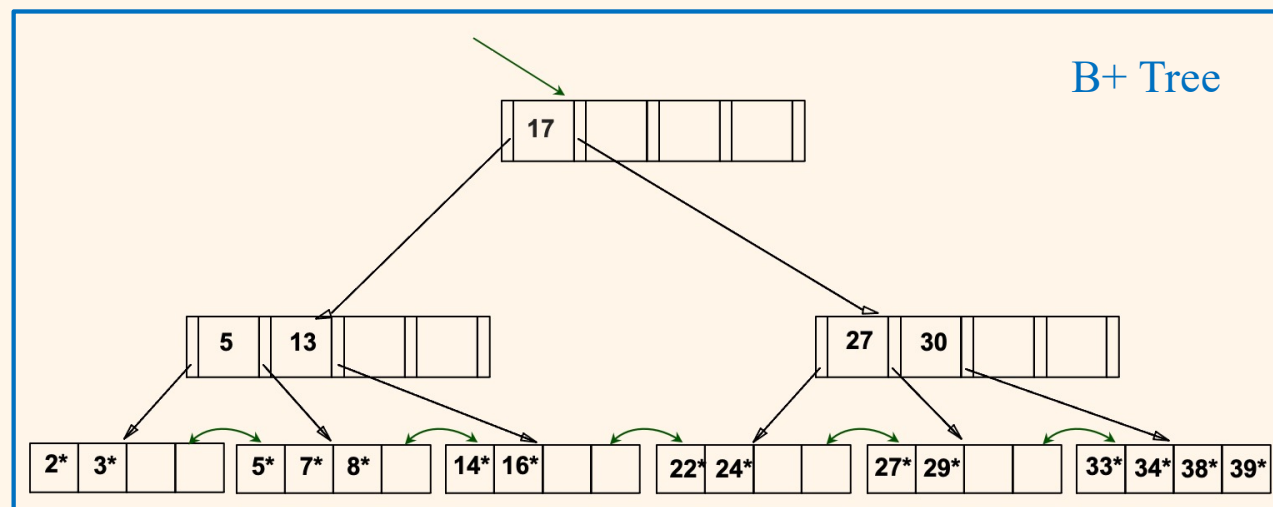
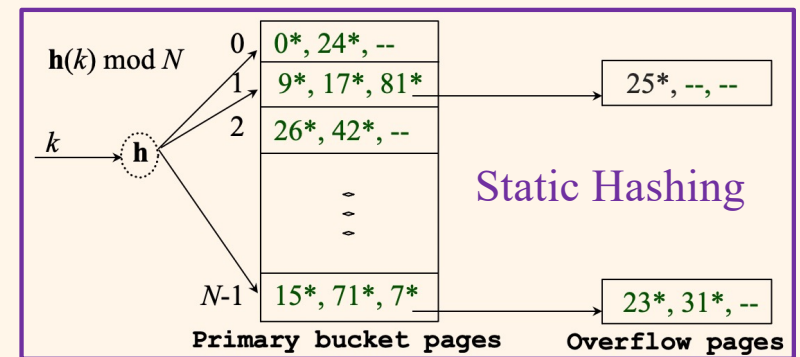
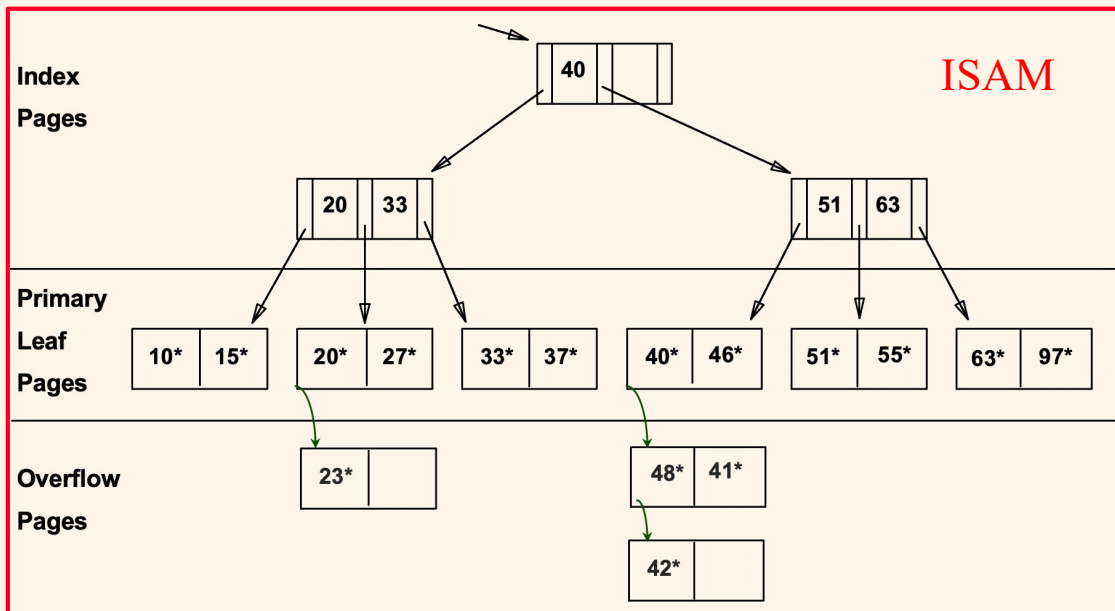
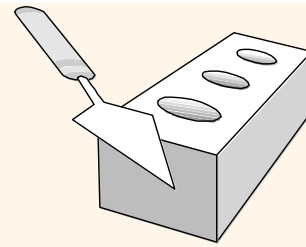
- ❖ Buckets contain *data entries* (like for ISAM or B+ trees) – very similar to what we just looked at.
- ❖ Hash function works on *search key* field of record *r*. Must distribute values over range  $0 \dots N-1$ .
  - $h(key) = (a * key + b) \bmod N$  works fairly well.
  - *a* and *b* are constants; lots known about how to tune **h**.
- ❖ **Long overflow chains** can develop and degrade performance. (Analogous to ISAM.)
  - **Extendible Hashing** and **Linear Hashing**: More dynamic approaches that address this problem. (Take CS122c!)



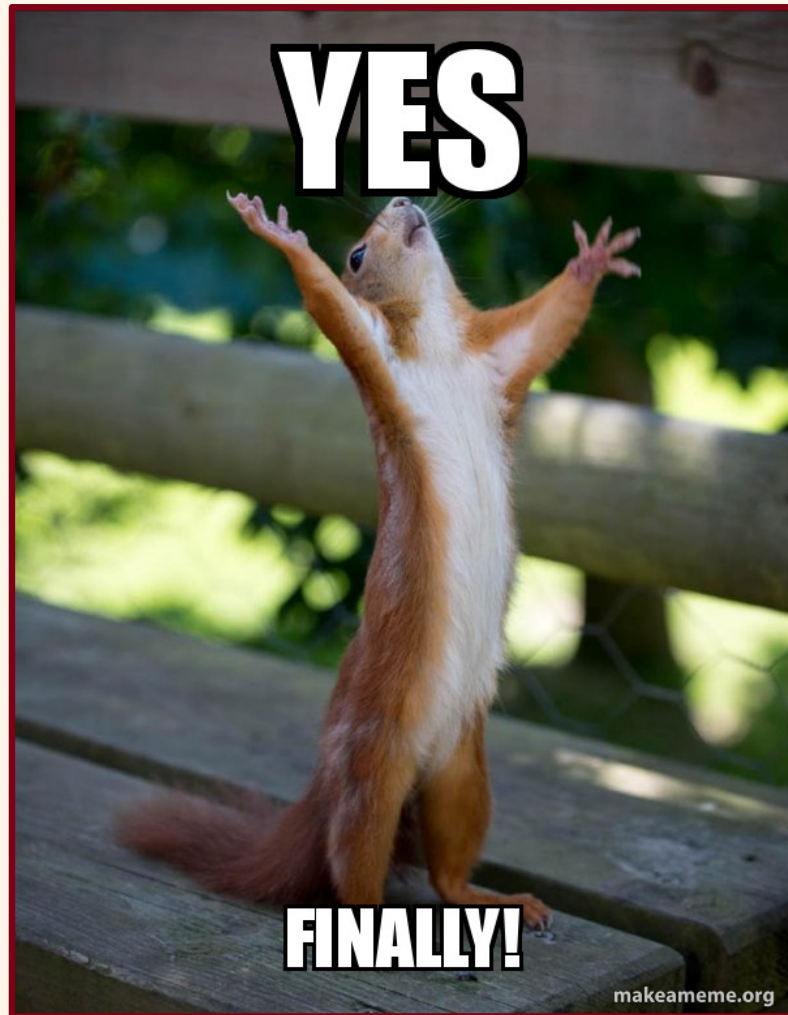
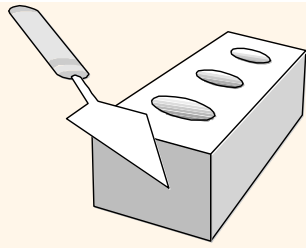
# *Indexing Summary*

- ❖ Tree-structured indexes are ideal for range-searches, also good for equality searches.
- ❖ **ISAM** is a static structure. (Prehistoric B+ Tree!)
  - Only leaf pages modified; overflow pages needed.
  - Overflow chains can degrade performance unless size of data set and data distribution stay constant.
- ❖ **B+ tree** is a dynamic structure. (Widely used!)
  - Inserts/deletes leave tree height-balanced;  $\log_F N$  cost.
  - High fanout  $F \rightarrow$  tree depth rarely more than 3-4.
- ❖ **Hashed indexes** are an option for equality searches.

# Indexing Summary (cont.)



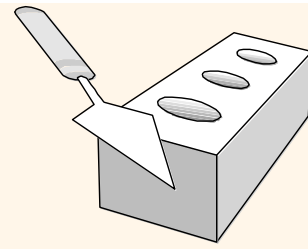
# *Next: Physical Database Design*





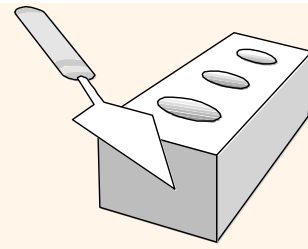
# *Physical DB Design: Overview*

- ❖ After ER design, schema refinement, and the definition of any views, we have the *conceptual* and *external* schemas for our database.
- ❖ Next step is to *choose indexes*, make *clustering decisions*, and *refine* the conceptual and external *schemas* (if needed) to meet *performance goals*.
- ❖ Start by *understanding the workload*:
  1. Most important queries and how often they arise.
  2. Most important updates and how often they arise.
  3. Desired performance goals for those queries/updates?



# *Decisions to Be Made Include...*

- ❖ What indexes should we create?
  - Which relations should have indexes? What field(s) should be their search keys? Should we build several indexes?
- ❖ For each index, what kind of an index should it be?
  - B+ tree? Hashed? Clustered? Unclustered?
- ❖ Should we make changes to the conceptual schema?
  - Consider alternative normalized schemas? (There are multiple choices when decomposing into BCNF, etc.)
  - Should we “undo” some decomposition steps and settle for a lower normal form? (*“Denormalization.”*)
  - Horizontal partitioning, materialized views, replication, ...



# *Understanding the Workload*

- ❖ For each **query** in the workload:
  - Which relations does it access?
  - Which attributes are retrieved?
  - Which attributes appear in selection/join conditions?  
(And *how selective* are those conditions expected to be?)
- ❖ For each **update** in the workload:
  - Which attributes are involved in selection/join conditions?  
(And *how selective* are those conditions likely to be?)
  - The type of update (INSERT/DELETE/UPDATE), and the attributes that are affected.



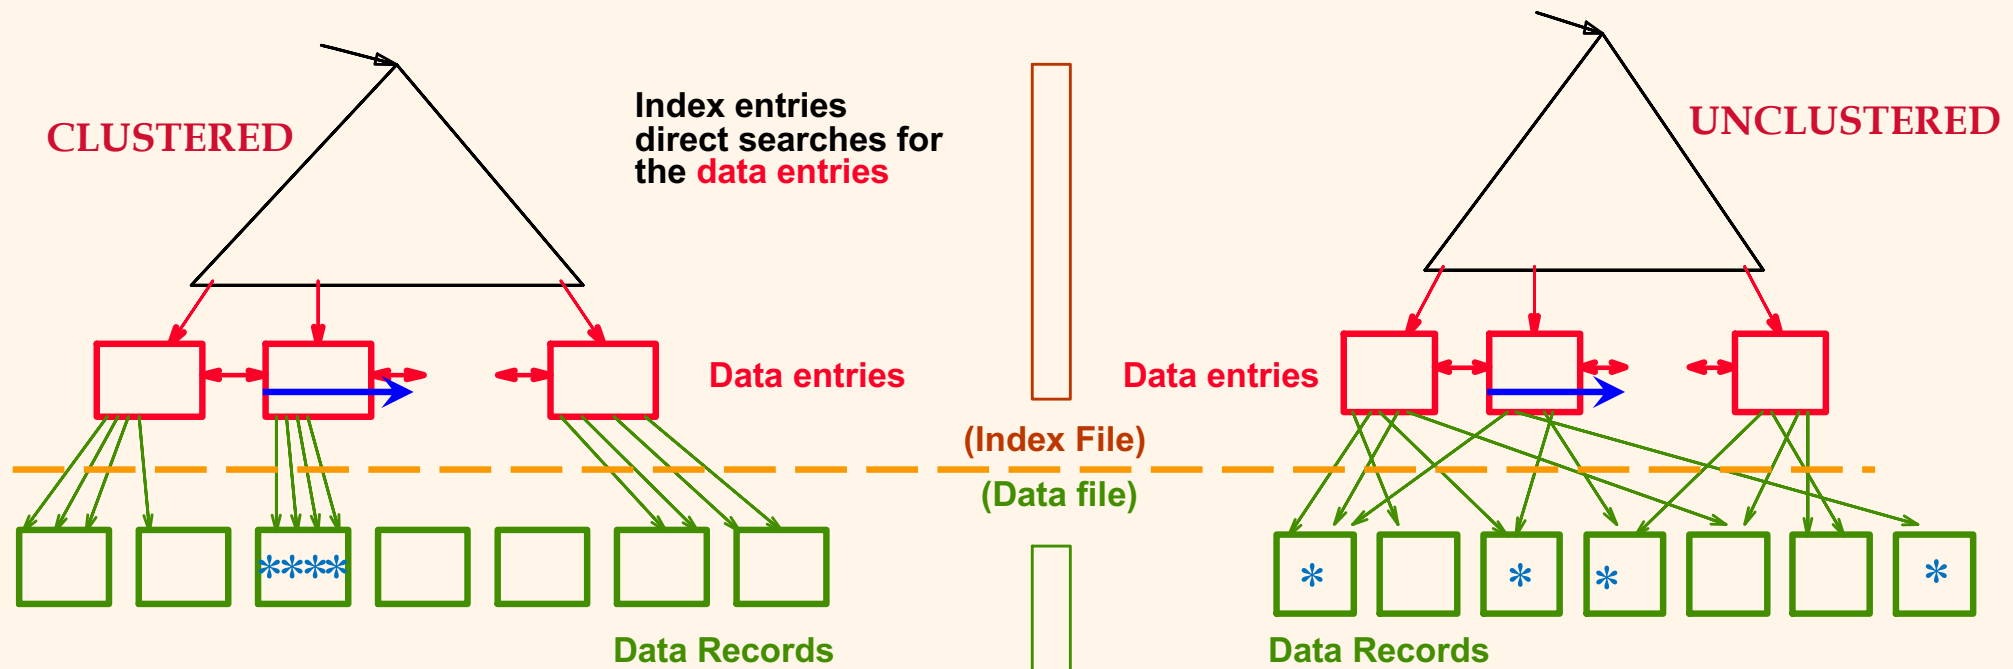
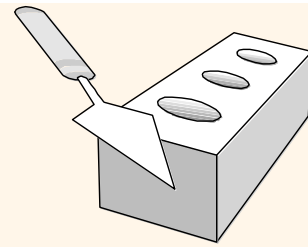


# *Index Classification (Review)*

- ❖ *Primary vs. secondary*: If index search key contains the primary key, this is called the *primary* index.
  - *Unique* index: Search key contains a *candidate* key.
- ❖ *Clustered vs. unclustered*: If the order of data entries is the same as, or nearly so, the order of stored data records, we have a clustered index.
  - A table can be clustered on *at most one* search key.
  - Cost of retrieving data records via an index varies *greatly* based on whether index is clustered or not!
  - Some systems always cluster on the primary key.

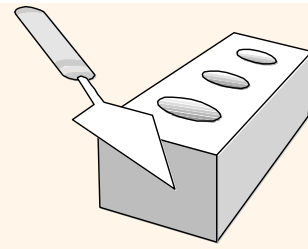


# Clustered vs. Unclustered Indexes (Reminder)



*(Read each page once.)*

*(Read more pages – and **repeatedly!**)*



## *Choice of Indexes (Cont'd.)*

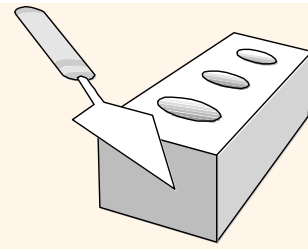
- ❖ **One approach:** Consider the most important queries in turn. Consider the best query plan using the current indexes, and see if a better plan is possible with an additional index. If so, create it.
  - This means we must understand and see how a DBMS evaluates its queries. (*Query execution plans.*)
  - Let's start by discussing simple 1-table queries!
- ❖ Before creating an index, must also consider its impact on updates in the workload.
  - *Trade-off:* Indexes can make queries go faster, but updates will become slower. (Indexes require disk space, too.)



# Index Selection Guidelines

- ❖ Attributes in **WHERE** clause are candidates for index keys.
  - Exact match condition → hashed index (or B+ tree).
  - Range query → B+ tree index.
    - Clustering especially useful for range queries, but can also help with *equality queries with **duplicate values*** (i.e., a non-key field index).
- ❖ **Multi-attribute** search keys should be considered when a WHERE clause contains several conditions.
  - Order of attributes in key matters for range queries.
  - Such indexes can sometimes enable **index-only** strategies for important queries (e.g., aggregates / grouped aggregates).
    - *Note:* For index-only strategies, data clustering isn't important!
- ❖ Choose indexes that benefit **as many queries** as possible.
  - Only **one** index can be clustered per relation, so choose it based on important queries that can benefit the most from clustering.

# Some Clustered Index Use Cases



```
SELECT E.dno
FROM   Emp E
WHERE  E.age > 40;
```

- ❖ B+ tree index on *E.age* can be used to get qualifying tuples.
  - How selective is the condition?
  - Should the index be clustered?

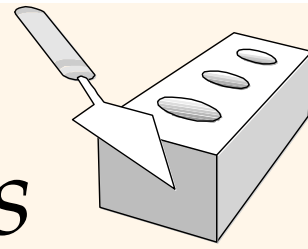
```
SELECT E.dno,
       COUNT (*)
FROM   Emp E
WHERE  E.age > 10
GROUP BY E.dno;
```

- ❖ Consider the GROUP BY query.
  - If most tuples have *E.age* > 10, using *E.age* index and grouping the retrieved tuples may be costly.
  - Clustered *E.dno* index may win!

```
SELECT E.dno
FROM   Emp E
WHERE  E.hobby='Stamps';
```

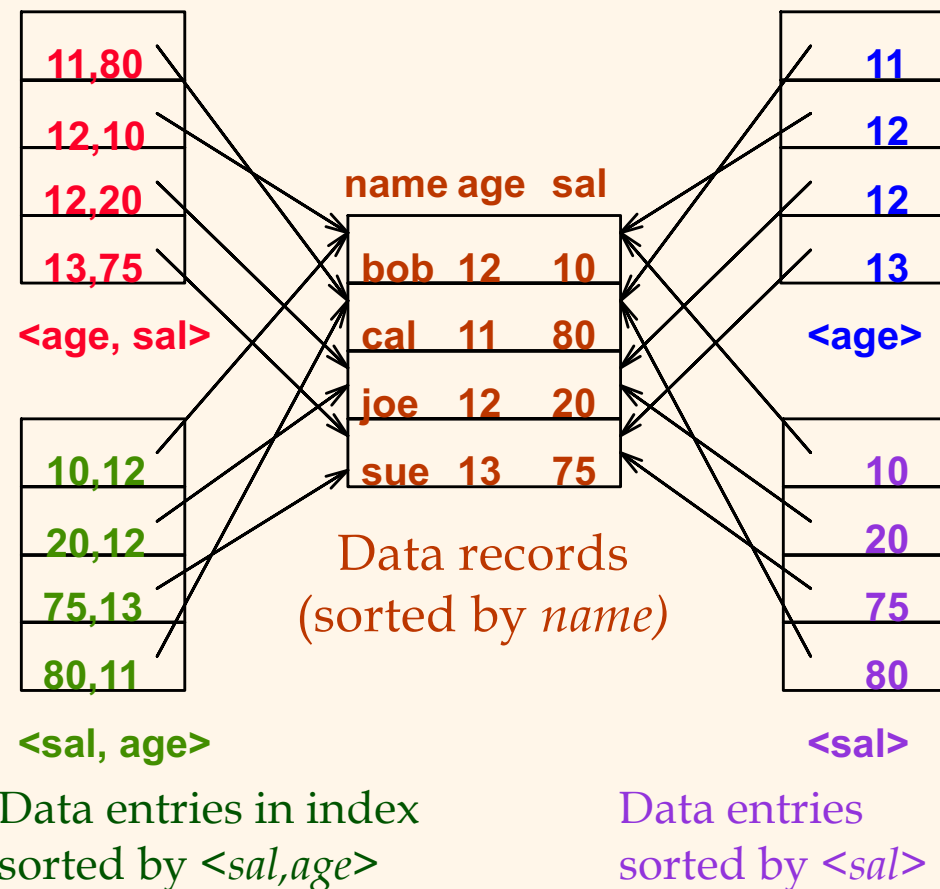
- ❖ Equality queries & duplicates:
  - Clustering on *E.hobby* helps!

# Indexes with Composite Search Keys

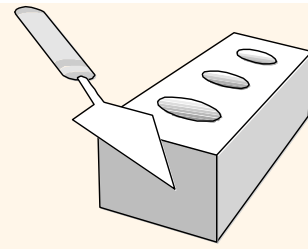


- ❖ **Composite Search Keys:** Search on a combination of fields.
  - **Equality query:** Every field value is equal to a constant value. E.g. wrt  $\langle \text{sal}, \text{age} \rangle$  index:
    - (age=20 AND sal=75)
  - **Range query:** Some field value is a range, not a constant. E.g. again wrt  $\langle \text{sal}, \text{age} \rangle$  index:
    - age=20; or (age=20 AND sal > 10)
- ❖ Data entries in index sorted by search key to support such range queries.
  - **Lexicographic order**

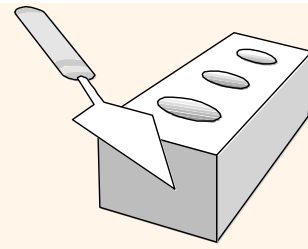
Various composite key indexes using lexicographic (ASC) order.



# Composite Search Keys



- ❖ To retrieve Emp records with  $age=30$  **AND**  $sal=4000$ , an index on  $\langle age, sal \rangle$  or  $\langle sal, age \rangle$  would be better than an index only on  $age$  or an index only on  $sal$ .
  - *Note*: Choice of index key is orthogonal to clustering.
- ❖ If condition is:  $20 < age < 30$  **AND**  $3000 < sal < 5000$ :
  - Clustered B+ tree index on  $\langle age, sal \rangle$  or  $\langle sal, age \rangle$  is best.
- ❖ If condition is:  $age=30$  **AND**  $3000 < sal < 5000$ :
  - Clustered  $\langle age, sal \rangle$  index *much* better than  $\langle sal, age \rangle$  index! (*Think about why*: Draw a picture of the index!)
- ❖ Composite indexes are larger; updated more often.



# Index-Only Query Plans

- ❖ Some queries can be answered without retrieving *any* tuples from one or more of the relations involved if a suitable index is available.

(Sometimes called a “*covering index*” for the given query.)

$\langle E.dno \rangle$   
 $\langle E.dno, E.sal \rangle$   
*B+ tree index!*

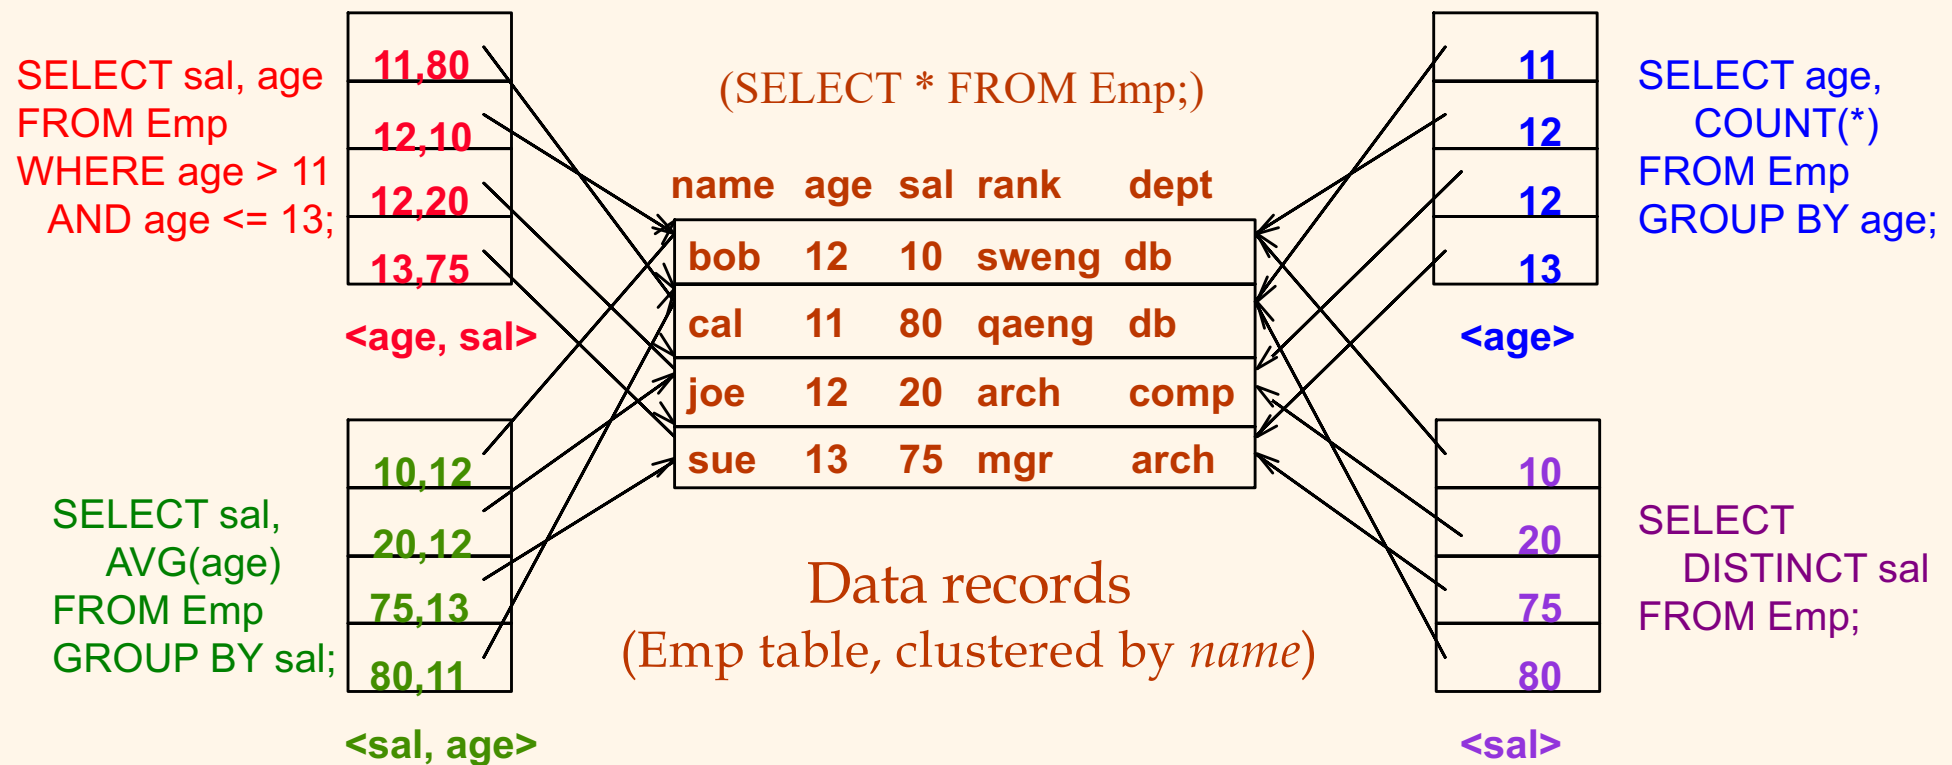
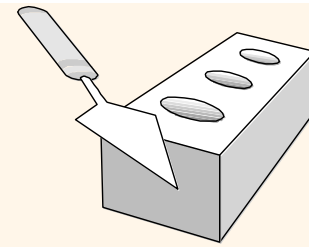
$\langle E.age, E.sal \rangle$   
*B+ tree index!*

```
SELECT E.dno, COUNT(*)  
FROM Emp E  
GROUP BY E.dno;
```

```
SELECT E.dno, MIN(E.sal)  
FROM Emp E  
GROUP BY E.dno;
```

```
SELECT AVG(E.sal)  
FROM Emp E  
WHERE E.age=25 AND  
E.sal BETWEEN 3000 AND 5000;
```

# Some Illustrated Index-Only Plans



Note: The index files are each much smaller than the main file!



*To Be Continued...*

