


Introduction to Data Management

**** The “Flipped” Edition ****




Lecture #21 (Physical DB Design II)

Instructor: Mike Carey
mjcarey@ics.uci.edu



Announcements




❖ Roadmap reminder:

Midterm Exam 2	Mon, Nov 15 (during lecture time)
Storage	Ch. 12.1-12.4, 12.6-12.7
Indexing	Ch. 14.1-14.4, 14.5
Physical DB Design	Ch. 14.6-14.7, 15.1-15.3, 15.5.3
Semistructured Data Management (<i>a.k.a.</i> NoSQL)	Ch. 8.1, ⇨ AsterixDB SQL++ Primer , ⇨ Couchbase SQL++ Book
Data Science 1: Advanced SQL Analytics	Ch. 5.5, 11.3
Data Science 2: Notebooks, Dataframes, and Python/Pandas	Lecture notes and Jupyter notebook
Basics of Transactions	Ch. 4.3, Ch. 17
Endterm Exam	Fri, Dec 3 (during lecture time)

- ❖ HW #6 should be wrapping up today by 6PM!
 - Solution will come out 24 hours after that, as usual
- ❖ Midterm #2 is now just a weekend away (*wow!*)
 - Monday (Nov. 15), in person
 - Gradescope + hard copy cheat sheet, assigned seats – you know the drill (☺)
- ❖ **Today:** Part II of II on Physical DB Design!

Join Methods & Index Selection




❖ When considering a join condition:

- Index Nested Loop join (**INLJ**) method:
 - For each outer table tuple, use its join attribute value to probe the inner table for tuples to join (match) it with.
 - **Indexing the inner table's join column** will help!
 - Good for this index to be *clustered* if the join column is *not* the inner's PK (e.g., FK) and inner tuples need to be fetched.
- Sort-Merge join (**SMJ**) method:
 - Sort outer and inner tables on join attribute value and then scan them concurrently to match tuples.
 - **Clustered B+ trees on both join column(s)** fantastic for this!
- Hash join (**HJ**) method:
 - Build in-memory hash table for one table, probe w/ the other.
 - Indexing not needed (not for the join part, anyway).

Example 1

```
SELECT E.ename, D.mgr  
FROM Emp E, Dept D  
WHERE D.dname = 'Toy' AND E.dno=D.dno;
```



- ❖ Index on *D.dname* supports the ‘Toy’ selection.
 - Given this, an index on *D.dno* is not needed (not used!).
- ❖ Index on *E.dno* allows us to then fetch the matching (inner) Emp tuples for each outer Dept tuple.
- ❖ What if **WHERE** included: “... AND *E.age*=25”?
 - Could instead retrieve Emp tuples using index on *E.age*, then join with Dept tuples also satisfying *dname* selection. (Similar to the strategy that used the *E.dno* index.)
 - So, if *E.age* index were already created, this particular query provides less support for adding an *E.dno* index.


Example 2

```
SELECT E.ename, D.mgr  
FROM Emp E, Dept D  
WHERE E.sal BETWEEN 10000 AND 20000  
AND E.hobby= 'Stamps' AND E.dno=D.dno;
```



- ❖ Clearly, Emp (E) should be the outer relation.
 - Suggests that we build an index (e.g., hashed) on *D.dno*.
- ❖ But what index should we build on Emp?
 - B+ tree on *E.sal* could be used, OR an index on *E.hobby* could be used. Only one of these is needed, and which is better depends upon the *selectivity* of the conditions.
 - As a *very* rough rule of thumb, equality selections tend to be more selective than range selections.
- ❖ As both examples indicate, our choice of indexes is guided by the plan(s) that we expect an optimizer to choose for a query. ∴ *Understand query optimizer!*


Clustering and Joins




```
SELECT E.ename, D.mgr  
FROM Emp E, Dept D  
WHERE D.dname= 'Toy' AND E.dno=D.dno;
```

- ❖ Clustering is very useful when accessing the inner tuples in an INLJ (index nested loops join).
 - Should make index on *E.dno* clustered. (Q: See why?)
- ❖ **Summary:** Clustering is useful whenever *many* tuples are expected for one value or for a range of values.

Tuning the Conceptual Schema




- ❖ The choice of conceptual schema should be guided by the workload, in addition to redundancy issues:
 - We may go for a 3NF (or lower!) schema rather than BCNF.
 - Workload may influence the choice we make when we're decomposing a relation into 3NF or BCNF.
 - We might *denormalize* (i.e., **undo** a decomposition step), or we might add redundant fields to a relation.
 - We might consider *vertical decompositions*.
- ❖ If such changes come *after* a database is in use, it's called *schema evolution*; might want to mask some of the changes from applications by defining *views*.



Some Example Schemas (& Tradeoffs)


Suppliers(sid, sname, address, phone, ...)
Parts(pid, pname, size, color, listprice, ...)
Stock(sid, pid, price, quantity)

- ❖ What if a large fraction of the workload consists of Stock queries that also want suppliers' names?
 - `SELECT s.sid, s.sname, AVG(t.price) FROM Suppliers s, Stock t WHERE s.sid = t.sid GROUP BY s.sid, s.sname;`
 - Consider: `ALTER TABLE Stock ADD COLUMN sname ...;`
 - This is denormalization (on purpose, for performance!)
 - If $sid \rightarrow sname$ and $sname \rightarrow sid$, Stock would then be in 3NF.
 - Q: If $sid \rightarrow sname$ (but not vice versa), what NF would it be in?



Vertical Partitioning

- ❖ Consider a table with many columns, not all of which are of interest to all queries.
 - Ex: Emp(eno, email, name, addr, salary, age, dno)
- ❖ A given workload might actually turn out to be a “union of sub-workloads” in reality.
 - Employee communications queries
 - Employee compensation queries/ analytics
 - Employee department queries/ analytics



Vertical Partitioning Example

eno	email	name	addr	salary	age	dno
1	joe@aol.com	Joe	1 Main St.	100000	25	10
2	sue@gmail.com	Sue	10 State St.	125000	28	20
3	zack@fb.com	Zack	100 Wall St.	2500000	40	30




(Vertical partitioning: ⋈)

eno	email	name	addr
1	joe@aol.com	Joe	1 Main St.
2	sue@gmail.com	Sue	10 State St.
3	zack@fb.com	Zack	100 Wall St.

eno	salary	age
1	100000	25
2	125000	28
3	2500000	40


eno	dno
1	10
2	20
3	30

(In the limit: We get a column store!)




Horizontal Partitioning

- ❖ Occasionally, we may want to instead replace a relation by a set of relations that are *selections*.
 - Each new relation has same schema (columns) as the original, but only a subset of the rows.
 - Collectively, the new relations contain all rows of the original. (Typically, the new relations are *disjoint*.)
 - The original relation is the UNION (ALL) of the new ones (i.e., rather than the JOIN of the new ones).



Potential Horizontal Rationale

- ❖ Suppose contracts with values over 10000 are subject to different rules. (This means queries on Contracts will frequently contain the condition *val* > 10000.)
- ❖ One approach to deal with this would be to create a clustered B+ tree index on Contracts(*val*).
- ❖ Another approach could be to replace Contracts by two relations, LargeContracts & SmallContracts, with the same attributes.
 - Clusters data like an index, but without index overhead.
 - Can then cluster on other (perhaps different!) attributes.



Horizontal Partitioning Example


eno	email	name	addr	salary	age	dno
1	joe@aol.com	Joe	1 Main St.	100000	25	10
2	sue@gmail.com	Sue	10 State St.	125000	28	20
3	zack@fb.com	Zack	100 Wall St.	2500000	40	30



(Horizontal partitioning: U)

eno	email	name	addr	salary	age	dno
1	joe@aol.com	Joe	1 Main St.	100000	25	10
2	sue@gmail.com	Sue	10 State St.	125000	28	20

eno	email	name	addr	salary	age	dno
3	zack@fb.com	Zack	100 Wall St.	2500000	40	30




Masking Schema Changes

```
CREATE VIEW Contracts(cid, sid, jid, did, pid, qty, val)
  AS SELECT * FROM LargeContracts
  UNION ALL
  SELECT * FROM SmallContracts
```


- ❖ Replacement of Contracts by LargeContracts and SmallContracts can be masked by this view.
- ❖ Note: queries with *val* > 10000 can be written against LargeContracts* for faster execution; users concerned with performance must be aware of this change.

(*The DBMS is unaware of the two tables' value constraints.)

In General: Tuning Up Your Queries




- ❖ If a query runs slower than expected, see if an index needs to be re-built, or if **table statistics** are too old.
- ❖ Sometimes, the DBMS may not be executing the plan you had in mind. Common areas of weakness:
 - Selections involving **arithmetic** or **LIKE expressions**.
 - Selections involving **OR** conditions.
 - Selections involving **null values**.
 - **Lack of advanced evaluation features** like some index-only strategies or certain join methods, or poor size estimation.
- ❖ Check the query plan!!! Then adjust the choice of indexes or maybe **rewrite the query (or view)**.




Miscellany for Query Tuning

- ❖ Minimize the use of **DISTINCT**: Don't use the D-word if duplicates are acceptable or if the answer contains a key.
- ❖ Consider the DBMS's use of indexes when writing arithmetic expressions: $E.age = 2 * D.age$ will benefit from an index on $E.age$, but it wouldn't benefit from an index on $D.age$!




Physical DB Design Summary

- ❖ End-to-end DB design consists of several tasks: *requirements analysis, conceptual design, schema refinement, physical design* and finally *tuning*.
 - In general, one goes back and forth between tasks to refine a DB design
 - Decisions made in one task can influence choices in another task.
- ❖ Understanding the *workload* for the application, and performance goals, is essential to good design.
 - What are the important queries and updates?
What attributes/relations are involved?



Summary (Cont'd.)

- ❖ The conceptual schema should perhaps be refined by considering performance criteria and workload:
 - May choose 3NF or a lower normal form over BCNF.
 - May choose among several alternative decompositions based on the expected workload.
 - May actually *denormalize*, or undo, some decompositions.
 - May consider further *vertical* (or in rare cases *horizontal*) decompositions.



Summary (Cont'd.)

- ❖ Over time, indexes may have to be fine-tuned (dropped, created, re-built, ...) for performance.
 - Be sure to examine the query plan(s) used by the system and adjust the choices of indexes appropriately.
- ❖ Sometimes the system may still not find a good plan:
 - Null values, arithmetic conditions, string expressions, use of ORs, etc., can “confuse” some query optimizers.
- ❖ So, may have to rewrite a particular query or view:
 - Might need to re-examine your complex nested queries, complex conditions, etc.