



Introduction to Data Management


*** *The “Flipped” Edition* ***

Lecture #26 (Transactions: The Final Frontier...!)

Instructor: Mike Carey
mjcarey@ics.uci.edu



Announcements

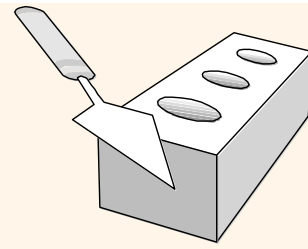


- ❖ HW wrap-up:
 - HW #8 is currently in flight!
 - Due **Wednesday** at **6 PM!** (Or Thursday at 6 PM if late)
- ❖ A note on Piazza points:
 - Pay attention to Piazza – don't miss the 1-point “final”! 😊
- ❖ Endterm exam: **Fri, Dec 3, in person, as usual!!!**
 - Same logistics as Midterms 1 & 2 (Gradescope + cheat sheet)
 - **Non**-cumulative (see Wiki syllabus for the official scope)
 - A sample exam and solution from the past are available
 - *Indexes, physical design, NoSQL/JSON, DS topics, transactions, ...*
- ❖ Final quiz and discussion sessions:
 - Quiz 10 will be available by Monday evening, and we'll also offer Endterm reviewing in this week's discussion sessions



Transactions

- ❖ Concurrent execution of user programs is essential for good DBMS performance (and wait times).
 - Disk I/Os are slow, so DBMSs keep the CPU cores busy by running multiple concurrent requests.
- ❖ A program may perform many operations on data from the DB, but the DBMS only cares about what's being read (R) and written (W) from/to the DB.
- ❖ A transaction is the DBMS's view of a user program:
 - It is seen as a sequence of database R's and W's.
 - The targets of the R's and W's are records (or pages).



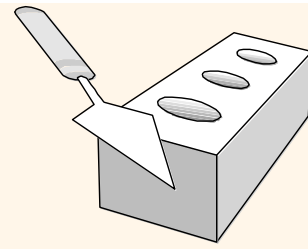
*The **ACID** Properties*

- ❖ Atomicity: Each transaction is **all or nothing**.
 - No worries about partial effects (if failures) and cleanup!
- ❖ Consistency: Each transaction moves the database from one **consistent state** to another one.
 - This is largely the application builder's responsibility...
- ❖ Isolation: Each transaction can be written as if it's the **only transaction** in existence.
 - No concurrency worries while building applications!
- ❖ Durability: Once a transaction has committed, its **effects will not be lost**.
 - Application code doesn't have to worry about data loss!



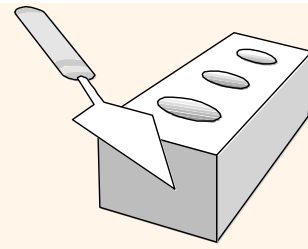
Concurrency in a DBMS

- ❖ Users run **transactions** and can think of each one as executing *all by itself*.
 - Concurrency is handled by the DBMS, which allows the actions (R's & W's) of various transactions to interleave.
 - Each transaction must leave the DB in a consistent state if it was consistent when the transaction started.
 - The DBMS may enforce some ICs, depending on the constraints declared in CREATE TABLE statements. (CHECK, PK, FK, ...)
 - But the DBMS does *not* understand the semantics of the data! (It doesn't know how interest on a bank account is computed.)
- ❖ Issues: Effects of *interleaving* and effects of *crashes*.



Atomicity of Transactions

- ❖ A transaction may *commit* after completing all of its actions, or it might *abort* (or might *be* aborted) after executing some of its actions.
 - Could violate a constraint, encounter some other error, be caught in a crash, or be picked to resolve a deadlock.
- ❖ The DBMS guarantees that transactions are *atomic*. A user can think of a Xact as doing **all** of its actions, in one step, or else executing **none** of its actions.
 - The DBMS *logs* all actions so that it can *undo* the actions of any aborted transactions.

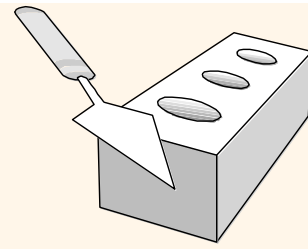


Example

- ❖ Consider two transactions (*Xacts*):

T1:	BEGIN	$A=A+100$,	$B=B-100$	END
T2:	BEGIN	$A=1.06*A$,	$B=1.06*B$	END

- ❖ E.g., T1 is transferring \$100 from bank account B to account A, while T2 is crediting both with 6% interest.
- ❖ No guarantee that T1 will execute before T2, or vice-versa, if both arrive together. The net effect *must* be *equivalent* to running them serially in some (either!) order.



A Quick Aside on “A” & “B”

❖ What *are* these two transactions, really?

```
T1: START TRANSACTION;  -- needed to couple the statements
    UPDATE Acct SET bal = bal + 100 WHERE acct_no = 101;
    UPDATE Acct SET bal = bal - 100 WHERE acct_no = 201;
    COMMIT;

T2: START TRANSACTION;  -- not needed if just one statement
    UPDATE Acct SET bal = bal * 1.06 WHERE acct_type = 'SV';
    COMMIT;
```

❖ Again, T1 is transferring \$100 from account B (201) to account A (101). T2 is giving all accounts their 6% interest payment.



Example (Cont'd.)

- ❖ Consider a possible interleaving (schedule):

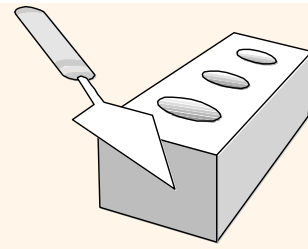
T1:	$A=A+100,$	$B=B-100$
T2:	$A=1.06*A,$	$B=1.06*B$

- ❖ That is OK. But what happens if:

T1:	$A=A+100,$	$B=B-100$
T2:	$A=1.06*A, B=1.06*B$	← <i>Too much interest!</i>


- ❖ The DBMSs view of the second schedule:

T1:	$R(A), W(A),$	$R(B), W(B)$
T2:	$R(A), W(A), R(B), W(B)$	



Scheduling Transactions (Defn's.)

- ❖ *Serial schedule*: Any schedule that does not interleave the actions of different transactions.
 - ❖ *Equivalent schedules*: If for any database state, the effect (on the DB) of executing the first schedule is identical to the effect of the second schedule.
 - ❖ *Serializable schedule*: A schedule that is equivalent to **some** (*any!*) serial execution of the transactions.
- ➔ If each transaction preserves consistency, then *every* serializable schedule preserves consistency!





Anomalies with Interleaved Execution

❖ Reading Uncommitted Data (“dirty reads”):



❖ Unrepeatable Reads:






Anomalies (Continued)

❖ Overwriting Uncommitted Data:

T7:	W(A),	W(B),	C
T8:	W(A),	W(B),	C

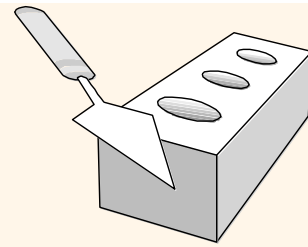
(Results are a “must have been concurrent!” mix of T7’s & T8’s writes – B from T7, and A from T8, yet both transactions wrote both A and B.)

Lock-Based Concurrency Control



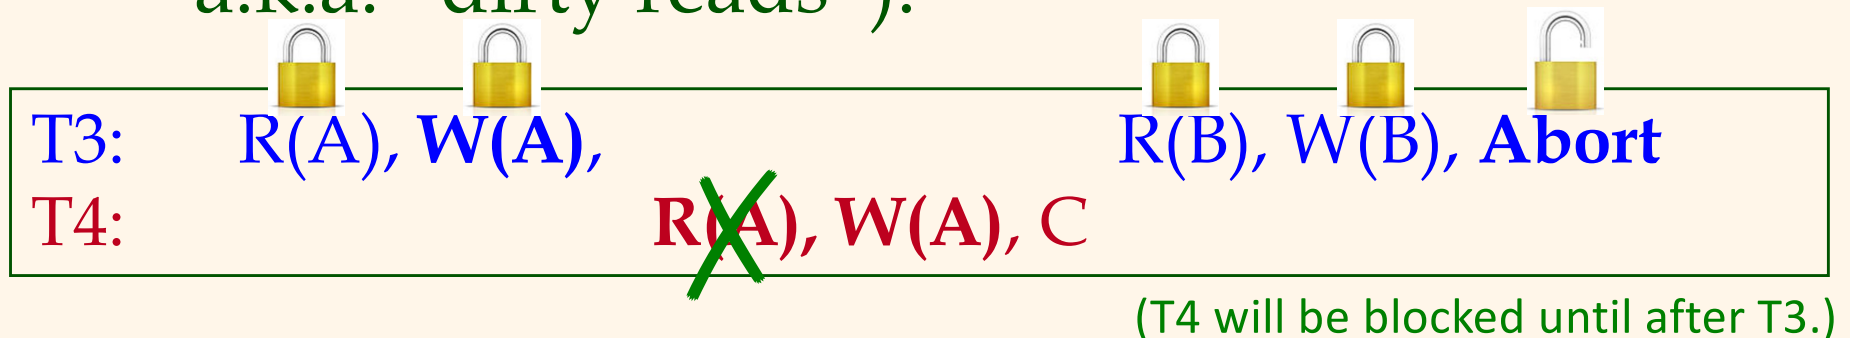
❖ Strict Two-phase Locking (2PL) Protocol:

- Each Xact acquires an **S (shared) lock** on an object before reading it, and an **X (exclusive) lock** on it before writing.
 - All of the locks held by a transaction are released only when the transaction completes.
 - *Note:* If a Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object – they must wait.
- ❖ Strict 2PL allows only serializable schedules.
- And additionally, it simplifies transaction aborts!

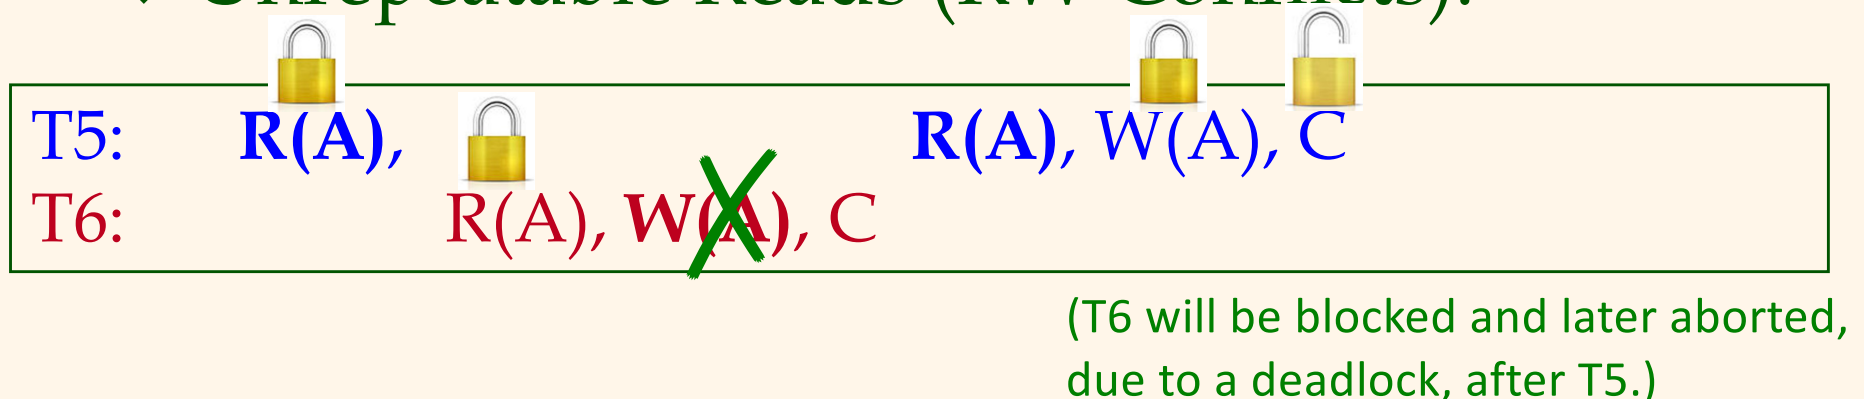



2PL Prevents the Anomalies

- ❖ Reading Uncommitted Data (WR Conflicts, a.k.a. “dirty reads”):




- ❖ Unrepeatable Reads (RW Conflicts):



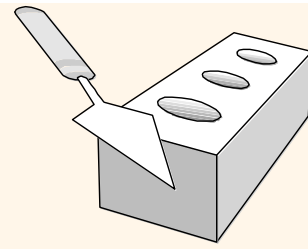


2PL & Anomalies (Continued)

❖ Overwriting Uncommitted Data (WW Conflicts):



(Now results will no longer be a “must have been concurrent!” intermingling of T7’s & T8’s writes...)



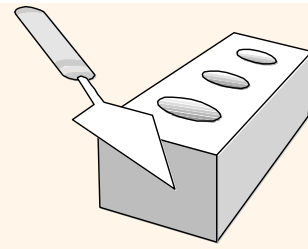
Aborting a Transaction

- ❖ If transaction T_i aborts, all its actions must be undone.
 - And, if some T_j already read a value last written by T_i , T_j must also be aborted! (“If I tell you, I’ll have to kill you...” 😊)
- ❖ Most systems avoid such *cascading aborts* by releasing a transaction’s locks only at *commit time*.
 - If T_i writes an object, T_j can read it only after T_i commits.
- ❖ In order to *undo* the actions of an aborted transaction, the DBMS keeps a *log* where every write is recorded.
 - Also used to recover from system crashes: active Xacts at crash time are aborted when the DBMS comes back up.




The Transaction Log

- ❖ The following actions are recorded in the log:
 - *Ti writes an object:* record its old and new values.
 - Log record must go to disk before the changed page – hence the name *write-ahead logging* (WAL).
 - *Ti commits/aborts:* write a log record noting the outcome.
- ❖ All log related activities (and all concurrency-related activities, like locking) are *transparently* taken care of by the DBMS.



Reminder: Disks and Files


- ❖ DBMSs store all information on disk.
- ❖ This has major implications for DBMS design!
 - **READ:** transfer data from disk to main memory (RAM).
 - **WRITE:** transfer data from RAM to disk.
 - Both are high-cost operations, relative to in-memory operations, so must be considered carefully!





Recovering From a Crash

- ❖ A three-phase recovery algorithm (*Aries*):
 - Analysis: Scan log (starting from most recent *checkpoint*) to identify the Xacts that were active, and the pages that were “dirty” in the buffer pool, when the system crashed.
 - Redo: Redo any updates to dirty pages to ensure that all logged updates were carried out and made it to disk.
(*Establishes the state from which to recover.*)
 - Undo: Undo the writes of all Xacts that were active at the crash (restoring the *before value* of each update from its log record), working backwards through the log, to abort any partially-completed transactions.



Support for Transactions in SQL-92

- ❖ A transaction is *automatically* started whenever a statement accesses or modifies the database
 - SELECT, UPDATE, CREATE TABLE, INSERT, ...
 - Multi-statement transactions also supported
- ❖ A transaction can be terminated by
 - A COMMIT statement
 - A ROLLBACK statement (SQL-speak for **abort**)
- ❖ Each transaction runs under a combination of an access mode and an isolation level



Transactions in SQL-92 (Cont'd.)

❖ Access mode – controls what the transaction can potentially do to the database:

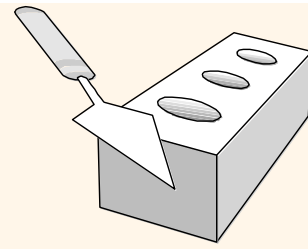
- READ ONLY: not permitted to modify the DB
- READ WRITE (*default*): allowed to modify the DB

❖ Isolation level – controls the transaction's exposure to other (concurrent) transactions:

- READ UNCOMMITTED: Can see “dirty” data!
- READ COMMITTED: Won't ever see dirty data.
- REPEATABLE READ: Re-reads get same result.
- SERIALIZABLE: No concurrency worries!

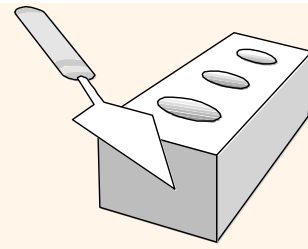


Increasing
isolation



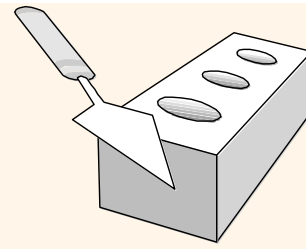
Which Isolation Level is for Me?

- ❖ An application-“controllable” tradeoff:
 - Consistency *vs.* performance (concurrency)
 - Warning: It will affect your programming model!
- ❖ Things to watch out for:
 - Default consistency level is DBMS engine-specific
 - Some engines may not support all levels
 - Default consistency level often not SERIALIZABLE
- ❖ You may also hear about “snapshot isolation”
 - DBMS keeps multiple versions of data
 - Transactions see versions as of when they started



Remember the **ACID** Properties!


- ❖ **Atomicity**: Each transaction is *all or nothing*.
 - No worries about partial effects (if failures) and cleanup.
- ❖ **Consistency**: Each transaction moves the database from one *consistent state* to another one.
 - This is largely the application builder's responsibility.
- ❖ **Isolation**: Each transaction can be written as if it's the *only transaction* in existence (*if so desired*).
 - Minimize concurrency worries when building applications.
- ❖ **Durability**: Once a transaction has committed, its *effects will not be lost*.
 - Application code needn't worry about data loss.



*A Few Quick **NoSQL** Xact Notes*

- ❖ For transactions, NoSQL systems tend to be limited to *record-level* transactions (in order to *scale* on a cluster)
- ❖ As a result, one sometimes considers an application's transactional needs when picking a schema (deciding what to "nest") for it, sigh!

You Made It!




Topic Coverage and Exam Schedule

Syllabus

Topic	Reading (Required!)
Databases and DB Systems	Ch. 1
Entity-Relationship (E-R) Data Model	Ch. 6.1-6.5, 6.8-6.9
Relational Data Model	Ch. 2.1-2.4, 3.1-3.2
E-R to Relational Translation	Ch. 6.6-6.7
Relational Design Theory	Ch. 7.1-7.4.2
Midterm Exam 1	Fri, Oct 22 (during lecture time)
Relational Algebra	Ch. 2.5-2.7
Relational Calculus	⇒ Wikipedia: Tuple relational calculus
SQL Basics (SPJ and Nested Queries)	Ch. 3.3-3.5
SQL Analytics: Aggregation, Nulls, and Outer Joins	Ch. 3.6-3.9, 4.1
Advanced SQL: Constraints, Triggers, Views, and Security	Ch. 4.2, 4.4-4.5, 4.7
Midterm Exam 2	Mon, Nov 15 (during lecture time)
Storage	Ch. 12.1-12.4, 12.6-12.7
Indexing	Ch. 14.1-14.4, 14.5
Physical DB Design	Ch. 14.6-14.7, 15.1-15.3, 15.5.3
Semistructured Data Management (<i>a.k.a.</i> NoSQL)	Ch. 8.1, ⇒ AsterixDB SQL++ Primer , ⇒ Couchbase SQL++ Book
Data Science 1: Advanced SQL Analytics	Ch. 5.5, 11.3
Data Science 2: Notebooks, Dataframes, and Python/Pandas	Lecture notes and Jupyter notebook
Basics of Transactions	Ch. 4.3, Ch. 17
Endterm Exam	Fri, Dec 3 (during lecture time)



“But Wait!.... I Need More....!!!”



- ❖ **CS122A** has just given you an “outside” view of database management systems.
- ❖ **CS122B** is available to give you a “programmer’s” view – with an emphasis on data-centric web applications.
- ❖ **CS122C** (*a.k.a.* CS222 lite) is available to give you an “insider’s” (engine developer’s) view of DB systems.
- ❖ **CS122D** goes “Beyond SQL Data Management” – NoSQL, Graph DBs, Hadoop/Spark, and more.
- ❖ **CS223** is available for learning all about transactions and distributed databases.
- ❖ **CS199** (independent project work) is also a possible avenue for gaining further experience.